

Team13_Uber&Lyft_Price_prediction

December 6, 2023

1 Team13 Uber and Lyft Surge Price Prediction

Team Members: I-An Chien (U25013252), Shuo Ming Kuo (U09774481), YI-Cheng Chung (U24763431), Yu Chin Chen(U11038917)

1.1 Background

As residents living in the Boston area, we benefit from a convenient public transportation system, including commuter rail, subway lines, and bus routes. However, despite these options, sometimes schedules or routes may not fit our specific needs, prompting us to rely on services like Uber and Lyft for convenience. These ride-sharing apps have become an integral part of transportation in the Boston area, and pricing is an important factor in the decision-making process. When considering whether to choose Uber or Lyft, the first thing that comes to mind is price, as people often open both apps to compare and choose the more economical option.

Several factors contribute to price fluctuations on these platforms. For example, demand and supply dynamics play a crucial role, with peak periods or high demand causing higher price. In addition, external factors such as weather conditions, traffic or distance can also affect the cost of a ride. The availability of different vehicle types and service levels also results in different price points.

1.2 Problem Statement

In this project, our primary objective is to provide a prediction of Uber and Lyft surge prices under different weather and geolocation factors, such as: temperature, rain, wind, pick up location, and destination, to enhance passengers' ability to make informed booking decisions.

We aim to solve the problem of fluctuating Uber and Lyft fares under the influence of unpredictable surge pricing, offering passengers a reliable tool to estimate the cost of their trips.

Our analysis is geared towards examining the Uber and Lyft trips in Boston during 11/26/2018 to 12/18/2018. By finding the most related factors that affect fare prices, we will first use Linear Regression, Decision Tree Regression, Elastic Net Regression, and XGBoost Regression models to find the most accurate model for prediction. Then, we will use feature importance to pinpoint the most important factors that influence price for Uber and Lyft separately.

1.3 Problem importance

Provide budgeting suggestions for passengers:

Our predictive model is designed to predict Uber and Lyft fare prices by understanding the factors that influence surge pricing, including interactions between factors. Our predictive model allows travelers to enter specific trip details such as pick up location, destination, and vehicle type. This feature provides price estimates to travelers, allowing them to plan their budget and travel effectively.

Provide trip planning suggestions for passengers:

For passengers who are more flexible on their departure time, our model can also provide alternative times with more affordable prices. Furthermore, for passengers who are unsure about their pick up location, they can use our model to discover suitable locations based on their budget.

1.4 Dataset Description

Accessing Kaggle's data, cab_rides, and weather datasets provides a valuable opportunity to delve into the dynamics of Uber and Lyft taxi prices during the week of November-December 2018. The cab_rides data set consists of 10 columns and 693,070 rows and provides a comprehensive view of the various types of taxis offered by Uber and Lyft, including pricing and location information.

The weather dataset has 8 columns and 6,275 rows and provides information about weather attributes for various locations and books. This includes temperature, precipitation (rain) and cloud cover information. Integrating weather data into the analysis can provide a better understanding of how external factors such as weather conditions affect Uber and Lyft prices.

Datasource : Uber & Lyft Cab prices on Kaggle: direct access link [here](#)

Number of rows: 693,071 rows

Number of columns:

- cab_rides: 10 columns
- weather: 8 columns

Variable types: string, number, date, integer, boolean

Data description for cab_rides

Column title	Description
distance	distance between source and destination
cab_type	Uber or Lyft
time_stamp	epoch time when data was queried
destination	destination of the ride
source	the starting point of the ride
price	price estimate for the ride in USD
surge_multiplier	the multiplier by which price was increased, default 1
id	unique identifier

Column title	Description
product_id	uber/lyft identifier for cab-type
name	Visible type of the cab eg: Uber Pool, UberXL

Data description for weather

Column title	Description
temp	Temperature in F
location	Location name
clouds	Clouds
pressure	pressure in mb
rain	rain in inches for the last hr
time_stamp	epoch time when row data was collected
humidity	thumidity in %
wind	wind speed in mph

2 Data Processing

2.1 Import library & Document

```
[ ]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from google.colab import drive

%matplotlib inline
import warnings
warnings.filterwarnings("ignore")
```

```
[ ]: drive.mount('/content/drive')
cab_df = pd.read_csv("/content/drive/MyDrive/BA_810_Group13/cab_rides.csv")
weather = pd.read_csv("/content/drive/MyDrive/BA_810_Group13/weather.csv")
```

Mounted at /content/drive

2.2 Data Cleaning

2.2.1 Aligning cab_rides and weather units

Since the cab_df and weather use different time frequencies, we added a new column “merge_date” to take the average weather data by hour, and matched it to the trips table on the hour extracted from the timestamp.

```
[ ]: # Unifing timestamp format and turn to datetime
cab_df['date_time'] = pd.to_datetime(cab_df['time_stamp']/1000, unit='s')
weather['date_time'] = pd.to_datetime(weather['time_stamp'], unit='s')
```

```
[ ]: # Create a new column for merging, and imputing rain columns with 0
cab_df['merge_date'] = cab_df['source'].astype(str) + " - " +
    ↳ cab_df['date_time'].dt.date.astype(str) + " - " + cab_df['date_time'].dt.
    ↳ hour.astype(str)
weather['merge_date'] = weather['location'].astype(str) + " - " +
    ↳ weather['date_time'].dt.date.astype(str) + " - " + weather['date_time'].dt.
    ↳ hour.astype(str)

groupby_value = weather.groupby(['merge_date']).mean().reset_index()
groupby_value['rain'].fillna(0,inplace=True)
```

```
[ ]: # Merge two dataframe and drop NAN
groupby_value.index = groupby_value['merge_date']
merged_df = cab_df.join(groupby_value,on=['merge_date'],rsuffix = '_w')
merged_df.dropna(inplace=True)
merged_df.head(5)
```

```
[ ]: distance cab_type time_stamp destination source price \
0 0.44 Lyft 1544952607890 North Station Haymarket Square 5.0
1 0.44 Lyft 1543284023677 North Station Haymarket Square 11.0
3 0.44 Lyft 1543553582749 North Station Haymarket Square 26.0
4 0.44 Lyft 1543463360223 North Station Haymarket Square 9.0
5 0.44 Lyft 1545071112138 North Station Haymarket Square 16.5
```

```
surge_multiplier id product_id \
0 1.0 424553bb-7174-41ea-aeb4-fe06d4f4b9d7 lyft_line
1 1.0 4bd23055-6827-41c6-b23b-3c491f24e74d lyft_premier
3 1.0 c2d88af2-d278-4bfd-a8d0-29ca77cc5512 lyft_luxsuv
4 1.0 e0126e1f-8ca9-4f2e-82b3-50505a09db9a lyft_plus
5 1.0 f6f6d7e4-3e18-4922-a5f5-181cdd3fa6f2 lyft_lux
```

```
name date_time \
0 Shared 2018-12-16 09:30:07.890000128
1 Lux 2018-11-27 02:00:23.676999936
3 Lux Black XL 2018-11-30 04:53:02.749000192
4 Lyft XL 2018-11-29 03:49:20.223000064
5 Lux Black 2018-12-17 18:25:12.138000128
```

```
merge_date merge_date_w \
0 Haymarket Square - 2018-12-16 - 9 Haymarket Square - 2018-12-16 - 9
1 Haymarket Square - 2018-11-27 - 2 Haymarket Square - 2018-11-27 - 2
3 Haymarket Square - 2018-11-30 - 4 Haymarket Square - 2018-11-30 - 4
4 Haymarket Square - 2018-11-29 - 3 Haymarket Square - 2018-11-29 - 3
5 Haymarket Square - 2018-12-17 - 18 Haymarket Square - 2018-12-17 - 18
```

```
temp clouds pressure rain time_stamp_w humidity wind
0 38.460 0.290000 1022.25 0.000 1.544954e+09 0.760000 7.68
```

1	44.065	0.995000	1002.88	0.106	1.543286e+09	0.895000	12.63
3	35.080	0.000000	1013.71	0.000	1.543554e+09	0.700000	5.25
4	37.680	0.433333	998.42	0.000	1.543461e+09	0.706667	11.16
5	40.780	0.930000	1000.15	0.000	1.545072e+09	0.790000	7.55

3 Descriptive Analysis

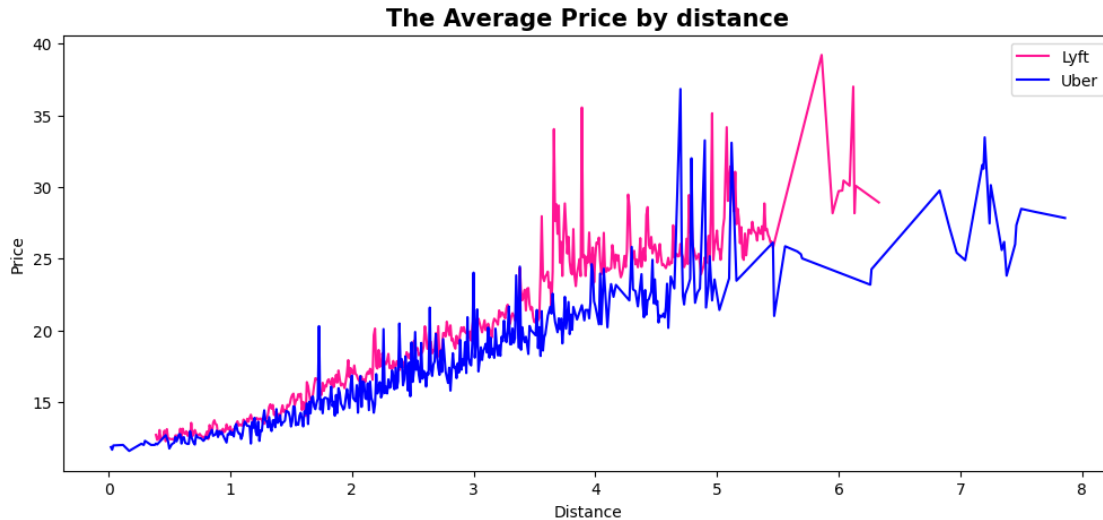
3.1 Which apps have lower prices, Uber or Lyft?

Customers usually want to save money, so they usually open both apps to compare prices before making a decision. One of the main factors that affects prices is the distance between origin and destination. Using this data, we can look at the average price for each app based on distance traveled.

```
[ ]: # plotting distance against price
fig , ax = plt.subplots(figsize = (12,5))
ax.plot(merged_df[merged_df['cab_type'] == 'Lyft'].groupby('distance').price.
        ↪mean().index,
        merged_df[merged_df['cab_type'] == 'Lyft'].groupby('distance').price.
        ↪mean(),
        label = 'Lyft', color='deeppink')

ax.plot(merged_df[merged_df['cab_type'] == 'Uber'].groupby('distance').price.
        ↪mean().index,
        merged_df[merged_df['cab_type'] == 'Uber'].groupby('distance').price.
        ↪mean(),
        label = 'Uber', color='blue')

ax.set_title('The Average Price by distance', fontsize= 15, fontweight='bold')
ax.set(xlabel = 'Distance', ylabel = 'Price' )
ax.legend()
plt.show()
```



From above figure, we can observe that for long-distance trips (distances exceeding 5 miles), Lyft tends to have a higher price. Therefore, it is advisable to opt for Uber in such cases.

However, for short-distance trips (distances less than 5 miles), Uber can occasionally have higher costs.

3.2 Do Uber and Lyft have different prices for various vehicle types?

Sometimes customers are traveling with friends or family and require a larger vehicle, and other times they seek luxury travel. Therefore, it becomes crucial to understand which apps customers should choose based on vehicle type.


```
[ ]: # plotting Vehicle Type against average price
uber_order = [ 'UberPool', 'WAV', 'UberX', 'UberXL', 'Black', 'Black SUV' ]
lyft_order = [ 'Shared', 'Lyft', 'Lyft XL', 'Lux', 'Lux Black', 'Lux Black XL' ]
fig, ax = plt.subplots(1, 2, figsize = (12, 4))
ax1 = sns.barplot(x = merged_df[merged_df['cab_type'] == 'Uber'].name,
                  y = merged_df[merged_df['cab_type'] == 'Uber'].price,
                  ax = ax[0], order = uber_order, palette='Blues')

ax2 = sns.barplot(x = merged_df[merged_df['cab_type'] == 'Lyft'].name,
                  y = merged_df[merged_df['cab_type'] == 'Lyft'].price,
                  ax = ax[1], order = lyft_order, palette='Reds')

ax1.set_xlabel = 'Vehicle Type', ylabel = 'Average Price')
ax2.set_xlabel = 'Vehicle Type', ylabel = 'Average Price')

ax1.set_title('The Uber Average Prices by Vehicle Type', fontweight='bold')
ax2.set_title('The Lyft Average Prices by Vehicle Type', fontweight='bold')
```

```
plt.show()
```



From above figures, comparing the price differences for various vehicle types, even when opting for shared rides, Uber (UberPool) tends to be more expensive than Lyft (Shared).

If you're looking for a smaller vehicle type, Lyft (Lyft) is generally more affordable compared to Uber (UberX).

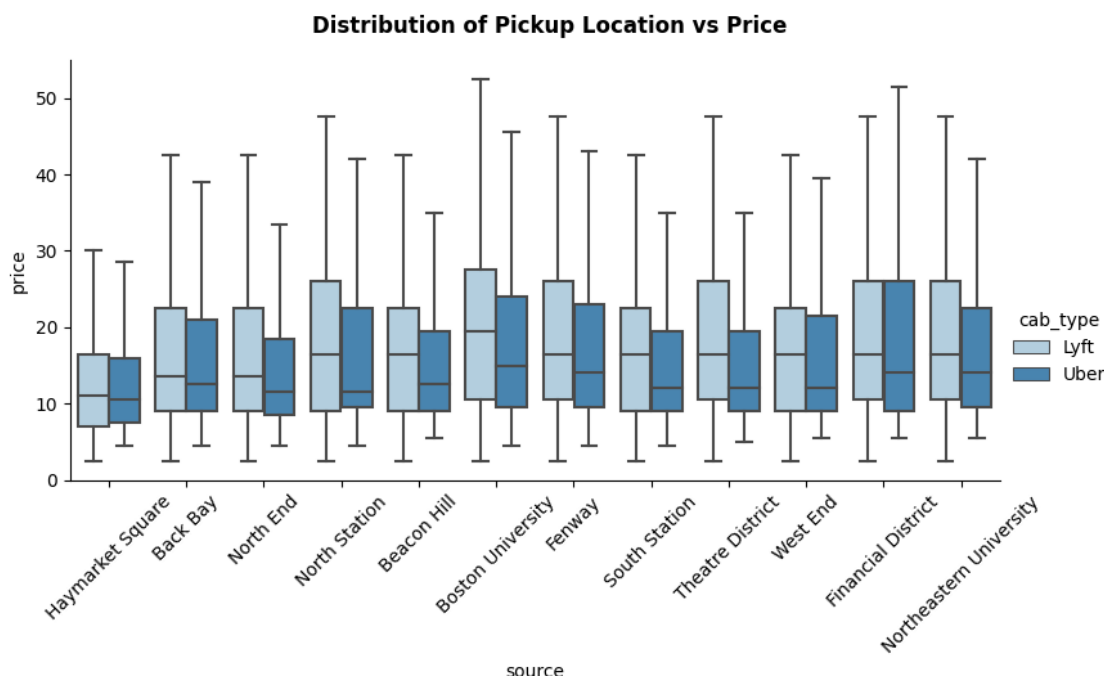
However, when considering larger vehicle types (XL), Uber (UberXL) typically comes with a higher price tag than Lyft (LyftXL).

For a premium service, it's recommended to choose Uber (UberBlack).

3.3 How do prices vary based on different pick-up locations?

It is possible that the pricing fluctuations are correlated with demand. Therefore, we would like to know whether there is a price difference among various pick-up locations. Understanding such variations can shed light on the impact of demand on pricing in different areas.

```
[ ]: # plotting source against price
sns.catplot(x='source', y='price', data=merged_df, kind='box', hue='cab_type',
            sym='', height=4, aspect=2, palette='Blues', dodge=True)
plt.tick_params(axis='x', rotation=45)
plt.suptitle("Distribution of Pickup Location vs Price", y=1.05,
            fontweight='bold')
plt.show()
```



From above figure, we are comparing the price differences from various pick up locations.

We observe that prices tend to be higher when starting from places like Boston University, Fenway, and Northeastern University.

One possible reason for this could be high demand in those areas. Additionally, the prices in the Financial District are more varied, which might be related to the distance of the trips.

4 Modeling

Which methods worked best for your problem?

We decided to implement three distinct search methods—GridSearch, RandomizedSearch, and Bayes Search—across four unique models—Linear Regression, Decision Tree Regression, XGBoost Regression and Elastic Net Regression—to achieve hyperparameter tuning. The goal was to identify the optimal models for both Uber and Lyft. To assess the models' performance, we compared the negative root mean squared error for each and selected the model with the highest value. Subsequently, we utilized the chosen model to make predictions on the test dataset. Notably, the best models for both datasets turned out to be XGBoost, utilizing the Bayes Search method.

Uber's best parameters in XGBoost:

1. colsample_bytree: 0.826784
2. max_depth: 5
3. n_estimators: 169

Lyft's best parameters in XGBoost:

1. colsample_bytree: 0.87526
2. max_depth: 7
3. n_estimators: 164

```
[ ]: # Selecting features
merged_df =
    merged_df[['distance', 'cab_type', 'destination', 'source', 'price', 'name', 'date_time', 'merge_d

# Creating time period variable and mapping
merged_df['hour'] = merged_df['date_time'].dt.hour.astype(str)
mapping = {
    '6': 'morning', '7' : 'morning', '8' : 'morning', '9' : 'morning',
    '10' : 'noon', '11' : 'noon', '12' : 'noon', '13' : 'noon',
    '14' : 'afternoon', '15' : 'afternoon', '16' : 'afternoon', '17' :
    'afternoon',
    '18' : 'evening', '19' : 'evening', '20' : 'evening', '21' : 'evening',
    '22' : 'night', '23' : 'night', '0' : 'night', '1' : 'night',
    '2' : 'night', '3' : 'late_night', '4' : 'late_night', '5' : 'late_night' }
merged_df['time_period'] = merged_df['hour'].replace(mapping)

merged_df.drop(columns=['date_time', 'merge_date', 'hour'], axis=1, inplace=True)
merged_df.head(5)
```

```
[ ]: distance cab_type destination source price name \
0 0.44 Lyft North Station Haymarket Square 5.0 Shared
1 0.44 Lyft North Station Haymarket Square 11.0 Lux
3 0.44 Lyft North Station Haymarket Square 26.0 Lux Black XL
4 0.44 Lyft North Station Haymarket Square 9.0 Lyft XL
5 0.44 Lyft North Station Haymarket Square 16.5 Lux Black

temp clouds pressure rain humidity wind time_period
0 38.460 0.290000 1022.25 0.000 0.760000 7.68 morning
1 44.065 0.995000 1002.88 0.106 0.895000 12.63 night
3 35.080 0.000000 1013.71 0.000 0.700000 5.25 late_night
4 37.680 0.433333 998.42 0.000 0.706667 11.16 late_night
5 40.780 0.930000 1000.15 0.000 0.790000 7.55 evening
```

```
[ ]: # Subsetting dataframe into uber and lyft
df_lyft = merged_df[merged_df['cab_type']=='Lyft'].copy()
df_uber = merged_df[merged_df['cab_type']=='Uber'].copy()

# Creating target and features
X_lyft = df_lyft.drop('price', axis=1)
y_lyft = df_lyft['price'].copy()

X_uber = df_uber.drop('price', axis=1)
y_yber = df_uber['price'].copy()
```

```
[ ]: # Info for Uber Dataframe
df_uber.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 329140 entries, 12 to 693070
Data columns (total 13 columns):
#   Column          Non-Null Count  Dtype
---  -
0   distance        329140 non-null float64
1   cab_type        329140 non-null object
2   destination     329140 non-null object
3   source          329140 non-null object
4   price           329140 non-null float64
5   name            329140 non-null object
6   temp            329140 non-null float64
7   clouds          329140 non-null float64
8   pressure        329140 non-null float64
9   rain            329140 non-null float64
10  humidity        329140 non-null float64
11  wind            329140 non-null float64
12  time_period     329140 non-null object
dtypes: float64(8), object(5)
memory usage: 35.2+ MB
```

```
[ ]: # Info for Lyft Dataframe
df_lyft.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 306102 entries, 0 to 693053
Data columns (total 13 columns):
#   Column          Non-Null Count  Dtype
---  -
0   distance        306102 non-null float64
1   cab_type        306102 non-null object
2   destination     306102 non-null object
3   source          306102 non-null object
4   price           306102 non-null float64
5   name            306102 non-null object
6   temp            306102 non-null float64
7   clouds          306102 non-null float64
8   pressure        306102 non-null float64
9   rain            306102 non-null float64
10  humidity        306102 non-null float64
11  wind            306102 non-null float64
12  time_period     306102 non-null object
dtypes: float64(8), object(5)
memory usage: 32.7+ MB
```

```
[ ]: # Splitting data into training and testing for Uber and Lyft dataframes
from sklearn.model_selection import train_test_split
X_train_lyft, X_test_lyft, y_train_lyft, y_test_lyft = train_test_split(X_lyft,
    ↪y_lyft, test_size=0.2, random_state=42)
X_train_uber, X_test_uber, y_train_uber, y_test_uber = train_test_split(X_uber,
    ↪y_yber, test_size=0.2, random_state=42)
```

```
[ ]: # Creating preprocessing pipeline
from sklearn import set_config
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import make_pipeline
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder

set_config(display='diagram')

cat_attribs = ["cab_type", "destination", "source", "name", "time_period"]
num_attribs = ["distance", "temp", "clouds", "pressure", "rain", "humidity",
    ↪"wind"]

preprocess_pipeline = ColumnTransformer([
    ("cat", OneHotEncoder(drop="first"), cat_attribs),
    ("num", StandardScaler(), num_attribs),])

preprocess_pipeline
```

```
[ ]: ColumnTransformer(transformers=[('cat', OneHotEncoder(drop='first'),
    ['cab_type', 'destination', 'source', 'name',
    'time_period']),
    ('num', StandardScaler(),
    ['distance', 'temp', 'clouds', 'pressure',
    'rain', 'humidity', 'wind'])])
```

```
[ ]: # Test if it's preprocessing
print(X_train_lyft.shape)
X_train_lyft_prepared = preprocess_pipeline.fit_transform(X_train_lyft)
print(X_train_lyft_prepared.shape)
```

```
(244881, 12)
(244881, 39)
```

```
[ ]: # Test if it's preprocessing
print(X_train_uber.shape)
X_train_uber_prepared = preprocess_pipeline.fit_transform(X_train_uber)
print(X_train_uber_prepared.shape)
```

```
(263312, 12)
```

(263312, 39)

4.1 Linear Regression

4.1.1 Uber

```
[ ]: # Linear regression for Uber dataframe
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import cross_val_score

lin_reg_uber = make_pipeline(preprocess_pipeline, LinearRegression())
lin_reg_uber.fit(X_train_uber, y_train_uber)
y_train_predictions_uber = lin_reg_uber.predict(X_train_uber)
lin_cv_rmse = -cross_val_score(lin_reg_uber, X_train_uber, y_train_uber,
                               scoring="neg_root_mean_squared_error", cv=3)
print(f"Average Linear Regression Cross-Validation RMSE: {lin_cv_rmse.mean():.
    0f}")
```

Average Linear Regression Cross-Validation RMSE: 2

```
[ ]: # Linear Regression Grid Search for Uber dataframe
from sklearn.model_selection import GridSearchCV
from sklearn.pipeline import Pipeline

param_grid = {
    'linearregression__fit_intercept': [True, False],
    'linearregression__copy_X': [True, False],
    'linearregression__positive': [True, False]
}

lin_uber_grid_search = GridSearchCV(lin_reg_uber, param_grid, cv=3,
    scoring='neg_root_mean_squared_error')
lin_uber_grid_search.fit(X_train_uber, y_train_uber)
```

```
[ ]: GridSearchCV(cv=3,
                estimator=Pipeline(steps=[('columntransformer',
                                           ColumnTransformer(transformers=[('cat',
OneHotEncoder(drop='first'),
['cab_type',
'destination',
'source',
'name',
'time_period'])],
                                           ('num',
StandardScaler(),
```

```

['distance',
 'temp',
 'clouds',
 'pressure',
 'rain',
 'humidity',
 'wind']]])),

        ('linearregression',
         LinearRegression()))],
param_grid={'linearregression__copy_X': [True, False],
            'linearregression__fit_intercept': [True, False],
            'linearregression__positive': [True, False]},
scoring='neg_root_mean_squared_error')

```

```

[ ]: # Presenting result
lin_uber_cv = pd.DataFrame(lin_uber_grid_search.cv_results_)
lin_uber_cv.sort_values(by="mean_test_score", ascending=False, inplace=True)
lin_uber_cv.filter(regex = '(^param_|mean_test_score)', axis=1)

```

```

[ ]:  param_linearregression__copy_X  param_linearregression__fit_intercept  \
1                                True                                True
5                                False                               True
3                                True                                False
7                                False                               False
0                                True                                True
2                                True                                False
4                                False                               True
6                                False                               False

    param_linearregression__positive  mean_test_score
1                                False        -2.402838
5                                False        -2.402838
3                                False        -2.402838
7                                False        -2.402838
0                                True             NaN
2                                True             NaN
4                                True             NaN
6                                True             NaN

```

```

[ ]: # Linear Regression Random Search for Uber dataframe

from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import randint

lin_uber_rnd_search = RandomizedSearchCV(
    lin_reg_uber, param_distributions=param_grid, n_iter=4, cv=3,

```

```

        scoring='neg_root_mean_squared_error', random_state=42)

lin_uber_rnd_search.fit(X_train_uber, y_train_uber)

```

```

[ ]: RandomizedSearchCV(cv=3,
                        estimator=Pipeline(steps=[('columntransformer',
ColumnTransformer(transformers=[('cat',
OneHotEncoder(drop='first'),
['cab_type',
'destination',
'source',
'name',
'time_period'])),
('num',
StandardScaler(),
['distance',
'temp',
'clouds',
'pressure',
'rain',
'humidity',
'wind'])])),
                        ('linearregression',
LinearRegression()))],
n_iter=4,
param_distributions={'linearregression__copy_X': [True,
False],
'linearregression__fit_intercept':
[True,
False],
'linearregression__positive': [True,
False]},
random_state=42, scoring='neg_root_mean_squared_error')

```

```

[ ]: # Presenting result
lin_uber_rnd_res = pd.DataFrame(lin_uber_rnd_search.cv_results_)
lin_uber_rnd_res.sort_values(by="mean_test_score", ascending=False,
↪ inplace=True)
lin_uber_rnd_res.filter(regex = '(^param_|mean_test_score)', axis=1)

```

```

[ ]:  param_linearregression__positive  param_linearregression__fit_intercept  \
0                                False                                True
1                                False                                True
3                                False                                False
2                                True                                 True

param_linearregression__copy_X  mean_test_score

```

0	True	-2.402838
1	False	-2.402838
3	False	-2.402838
2	True	NaN

```
[ ]: # Linear Regression Bayes Search for Uber dataframe
```

```
!pip install scikit-optimize
from sklearn.linear_model import LinearRegression
from skopt import BayesSearchCV
from skopt.space import Categorical
```

Collecting scikit-optimize

Downloading scikit_optimize-0.9.0-py2.py3-none-any.whl (100 kB)
100.3/100.3

kB 2.4 MB/s eta 0:00:00

Requirement already satisfied: joblib>=0.11 in

/usr/local/lib/python3.10/dist-packages (from scikit-optimize) (1.3.2)

Collecting pyaml>=16.9 (from scikit-optimize)

Downloading pyaml-23.9.7-py3-none-any.whl (23 kB)

Requirement already satisfied: numpy>=1.13.3 in /usr/local/lib/python3.10/dist-packages (from scikit-optimize) (1.23.5)

Requirement already satisfied: scipy>=0.19.1 in /usr/local/lib/python3.10/dist-packages (from scikit-optimize) (1.11.4)

Requirement already satisfied: scikit-learn>=0.20.0 in

/usr/local/lib/python3.10/dist-packages (from scikit-optimize) (1.2.2)

Requirement already satisfied: PyYAML in /usr/local/lib/python3.10/dist-packages (from pyaml>=16.9->scikit-optimize) (6.0.1)

Requirement already satisfied: threadpoolctl>=2.0.0 in

/usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.20.0->scikit-optimize) (3.2.0)

Installing collected packages: pyaml, scikit-optimize

Successfully installed pyaml-23.9.7 scikit-optimize-0.9.0

```
[ ]: sparse_columns = ['cab_type', 'destination', 'source', 'name', 'time_period']
```

```
# One-hot encoding categorical features without creating sparse matrices
```

```
X_train_uber_selected = X_train_uber[sparse_columns].to_numpy()
```

```
uber_bayes_encoder = OneHotEncoder(sparse=False)
```

```
X_train_uber_encoded = pd.get_dummies(X_train_uber[sparse_columns])
```

```
# Concatenating encoded features with the original dataset
```

```
X_train_uber_bayes_final = pd.concat([X_train_uber.
```

```
↳drop(columns=sparse_columns), X_train_uber_encoded], axis=1)
```

```
param_space = {
```

```
    'fit_intercept': Categorical([True, False]),
```

```

    'copy_X': Categorical([True, False]),
    'positive': Categorical([True, False])
}

lin_uber_bayesian_search = BayesSearchCV(
    estimator=LinearRegression(),
    search_spaces=param_space,
    n_iter=8,
    cv=3,
    scoring='neg_root_mean_squared_error',
    random_state=42
)

lin_uber_bayesian_search.fit(X_train_uber_bayes_final, y_train_uber)

```

```

[ ]: BayesSearchCV(cv=3, estimator=LinearRegression(), n_iter=8, random_state=42,
                  scoring='neg_root_mean_squared_error',
                  search_spaces={'copy_X': Categorical(categories=(True, False),
prior=None),
                                'fit_intercept': Categorical(categories=(True,
False), prior=None),
                                'positive': Categorical(categories=(True, False),
prior=None)})

```

```

[ ]: # Presenting result
lin_uber_bayes_res = pd.DataFrame(lin_uber_bayesian_search.cv_results_)
lin_uber_bayes_res.sort_values(by="mean_test_score", ascending=False,
                               inplace=True)
lin_uber_bayes_res.filter(regex = '(\param_|mean_test_score)', axis=1)

```

```

[ ]:  param_copy_X  param_fit_intercept  param_positive  mean_test_score
3             True                False             True      -2.402829
4             True                False             True      -2.402829
1             True                 True             False      -2.402838
2             False               True             False      -2.402838
5             True                 True             False      -2.402838
6             True                 True             False      -2.402838
7             True                 True             False      -2.402838
0             False               True              True      -2.402908

```

4.1.2 Lyft

```

[ ]: # Linear regression for Lyft dataframe

lin_reg_lyft = make_pipeline(preprocess_pipeline, LinearRegression())
lin_reg_lyft.fit(X_train_lyft, y_train_lyft)
y_train_predictions_lyft = lin_reg_lyft.predict(X_train_lyft)

```



```
lin_cv_rmse = -cross_val_score(lin_reg_lyft, X_train_lyft, y_train_lyft,
                                scoring="neg_root_mean_squared_error", cv=3)
print(f"Average Linear Regression Cross-Validation RMSE: {lin_cv_rmse.mean():.
    0f}")
```

Average Linear Regression Cross-Validation RMSE: 3

```
[ ]: # Linear Regression Grid Search for Lyft dataframe
```

```
param_grid = {
    'linearregression__fit_intercept': [True, False],
    'linearregression__copy_X': [True, False],
    'linearregression__positive': [True, False]
}

lin_lyft_grid_search = GridSearchCV(lin_reg_lyft, param_grid, cv=3,
    scoring='neg_root_mean_squared_error')
lin_lyft_grid_search.fit(X_train_lyft, y_train_lyft)
```

```
[ ]: GridSearchCV(cv=3,
                  estimator=Pipeline(steps=[('columntransformer',
                                             ColumnTransformer(transformers=[('cat',
OneHotEncoder(drop='first'),
['cab_type',
'destination',
'source',
'name',
'time_period'])),
                                             ('num',
StandardScaler(),
['distance',
'temp',
'clouds',
'pressure',
'rain',
'humidity',
'wind'])])),
                  ('linearregression',
                   LinearRegression()))],
                  param_grid={'linearregression__copy_X': [True, False],
                              'linearregression__fit_intercept': [True, False],
                              'linearregression__positive': [True, False]},
                  scoring='neg_root_mean_squared_error')
```

```
[ ]: # Presenting result
```

```
lin_lyft_cv = pd.DataFrame(lin_lyft_grid_search.cv_results_)
lin_lyft_cv.sort_values(by="mean_test_score", ascending=False, inplace=True)
```

```
lin_lyft_cv.filter(regex = '(^param_|mean_test_score)', axis=1)
```

```
[ ]: param_linearregression__copy_X param_linearregression__fit_intercept \
1          True          True
5          False         True
3          True          False
7          False         False
0          True          True
2          True          False
4          False         True
6          False         False
```

```
param_linearregression__positive mean_test_score
1          False        -3.47924
5          False        -3.47924
3          False        -3.47924
7          False        -3.47924
0          True          NaN
2          True          NaN
4          True          NaN
6          True          NaN
```

```
[ ]: # Linear Regression Random Search for Lyft dataframe

from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import randint

lin_lyft_rnd_search = RandomizedSearchCV(
    lin_reg_lyft, param_distributions=param_grid, n_iter=4, cv=3,
    scoring='neg_root_mean_squared_error', random_state=42)

lin_lyft_rnd_search.fit(X_train_lyft, y_train_lyft)
```

```
[ ]: RandomizedSearchCV(cv=3,
                        estimator=Pipeline(steps=[('columntransformer',
ColumnTransformer(transformers=[('cat',
OneHotEncoder(drop='first'),
['cab_type',
'destination',
'source',
'name',
'time_period'])),
('num',
StandardScaler(),
['distance',
'temp',
```

```

'clouds',
'pressure',
'rain',
'humidity',
'wind']]])),

                                ('linearregression',
                                LinearRegression()))],

n_iter=4,
param_distributions={'linearregression__copy_X': [True,
                                                    False],
                    'linearregression__fit_intercept':
[True,
False],

                    'linearregression__positive': [True,
                                                    False]},

random_state=42, scoring='neg_root_mean_squared_error')

```

```

[ ]: # Presenting result
lin_lyft_rnd_res = pd.DataFrame(lin_lyft_rnd_search.cv_results_)
lin_lyft_rnd_res.sort_values(by="mean_test_score", ascending=False,
                              inplace=True)
lin_lyft_rnd_res.filter(regex = '(^param_|mean_test_score)', axis=1)

```

```

[ ]:  param_linearregression__positive  param_linearregression__fit_intercept  \
0                                False                                True
1                                False                                True
3                                False                                False
2                                True                                 True

    param_linearregression__copy_X  mean_test_score
0                                True      -3.47924
1                                False     -3.47924
3                                False     -3.47924
2                                True              NaN

```

```

[ ]: # Linear Regression Bayes Search for Lyft dataframe

X_train_lyft_selected = X_train_lyft[sparse_columns].to_numpy()
lyft_bayes_encoder = OneHotEncoder(sparse=False)
X_train_lyft_encoded = pd.get_dummies(X_train_lyft[sparse_columns])
X_train_lyft_bayes_final = pd.concat([X_train_lyft.
                                       drop(columns=sparse_columns), X_train_lyft_encoded], axis=1)

param_space = {
    'fit_intercept': Categorical([True, False]),
    'copy_X': Categorical([True, False]),

```

```

        'positive': Categorical([True, False])
    }

    lin_lyft_bayesian_search = BayesSearchCV(
        estimator=LinearRegression(),
        search_spaces=param_space,
        n_iter=8,
        cv=3,
        scoring='neg_root_mean_squared_error',
        random_state=42
    )

    lin_lyft_bayesian_search.fit(X_train_lyft_bayes_final, y_train_lyft)

```

```

[ ]: BayesSearchCV(cv=3, estimator=LinearRegression(), n_iter=8, random_state=42,
                  scoring='neg_root_mean_squared_error',
                  search_spaces={'copy_X': Categorical(categories=(True, False),
prior=None),
                                'fit_intercept': Categorical(categories=(True,
False), prior=None),
                                'positive': Categorical(categories=(True, False),
prior=None)})

```

```

[ ]: # Presenting result
lin_lyft_bayes_res = pd.DataFrame(lin_lyft_bayesian_search.cv_results_)
lin_lyft_bayes_res.sort_values(by="mean_test_score", ascending=False,
                               inplace=True)
lin_lyft_bayes_res.filter(regex = '(~param_|mean_test_score)', axis=1)

```

```

[ ]:
param_copy_X param_fit_intercept param_positive mean_test_score
1          True                True          False      -3.479240
2          False                True          False      -3.479240
5          True                True          False      -3.479240
6          True                True          False      -3.479240
7          True                True          False      -3.479240
0          False                True           True      -3.479268
3          True                False           True      -3.523841
4          True                False           True      -3.523841

```

4.2 DecisionTree Regression

4.2.1 Uber

```

[ ]: # DecisionTreeRegressor for Uber dataframe
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import cross_val_score

```

```

from sklearn.pipeline import Pipeline

tree_uber = make_pipeline(preprocess_pipeline,
    ↳DecisionTreeRegressor(random_state=42))

tree_uber.fit(X_train_uber, y_train_uber)
tree_cv_rmsses = -cross_val_score(tree_uber, X_train_uber, y_train_uber,
    scoring="neg_root_mean_squared_error", cv=3)
print(f"Average DecisionTree Regression Cross-Validation RMSE: {tree_cv_rmsses.
    ↳mean():.0f}")

```

Average DecisionTree Regression Cross-Validation RMSE: 3

```

[ ]: # DecisionTreeRegressor Grid Search for Uber dataframe
from sklearn.model_selection import GridSearchCV
import numpy as np

param_grid = {
    'decisiontreeregressor__max_depth' : [10,20],
    'decisiontreeregressor__min_samples_split' : [2,5,10],
    'decisiontreeregressor__min_samples_leaf' : [1,5,10]
}

tree_uber_grid_search = GridSearchCV(tree_uber, param_grid, cv=3,
    ↳scoring='neg_root_mean_squared_error')
tree_uber_grid_search.fit(X_train_uber, y_train_uber)

```

```

[ ]: GridSearchCV(cv=3,
    estimator=Pipeline(steps=[('columntransformer',
    ColumnTransformer(transformers=[('cat',
OneHotEncoder(drop='first'),
['cab_type',
'destination',
'source',
'name',
'time_period'])),
('num',
StandardScaler(),
['distance',
'temp',
'clouds',
'pressure',
'rain',
'humidity',
'wind'])])),
    ('decisiontreeregressor',
DecisionTreeRegressor(random_state=42))]),

```

```

param_grid={'decisiontreeregressor__max_depth': [10, 20],
            'decisiontreeregressor__min_samples_leaf': [1, 5, 10],
            'decisiontreeregressor__min_samples_split': [2, 5, 10]},

scoring='neg_root_mean_squared_error')

```

```

[ ]: # Presenting result
tree_uber_cv = pd.DataFrame(tree_uber_grid_search.cv_results_)
tree_uber_cv.sort_values(by="mean_test_score", ascending=False, inplace=True)
tree_uber_cv.filter(regex = '^param_|mean_test_score', axis=1).head(10)

```

```

[ ]:  param_decisiontreeregressor__max_depth  \
6                                           10
7                                           10
8                                           10
3                                           10
4                                           10
5                                           10
2                                           10
1                                           10
0                                           10
15                                          20

    param_decisiontreeregressor__min_samples_leaf  \
6                                           10
7                                           10
8                                           10
3                                           5
4                                           5
5                                           5
2                                           1
1                                           1
0                                           1
15                                          10

    param_decisiontreeregressor__min_samples_split  mean_test_score
6                                           2      -1.935140
7                                           5      -1.935140
8                                          10      -1.935140
3                                           2      -1.938275
4                                           5      -1.938275
5                                          10      -1.938275
2                                          10      -1.944237
1                                           5      -1.946858
0                                           2      -1.958493
15                                          2      -2.023783

```

```
[ ]: # DecisionTreeRegressor Random Search for Uber dataframe
from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import randint

param_grid = {
    'decisiontreeregressor__max_depth' : randint(10,20),
    'decisiontreeregressor__min_samples_split' : randint(2,20),
    'decisiontreeregressor__min_samples_leaf' : randint(1,20)
}

tree_uber_rnd_search = RandomizedSearchCV(
    tree_uber, param_grid, n_iter=10, cv=3,
    scoring='neg_root_mean_squared_error', random_state=42)

tree_uber_rnd_search.fit(X_train_uber, y_train_uber)

[ ]: RandomizedSearchCV(cv=3,
                        estimator=Pipeline(steps=[('columntransformer',
ColumnTransformer(transformers=[('cat',
OneHotEncoder(drop='first'),
['cab_type',
'destination',
'source',
'name',
'time_period'])),
('num',
StandardScaler(),
['distance',
'temp',
'clouds',
'pressure',
'rain',
'humidity',
'wind'])])),
                        ('decisiontreeregressor',
DecisionTreeRegressor(random_state=4...
                        param_distributions={'decisiontreeregressor__max_depth':
<scipy.stats._distn_infrastructure.rv_discrete_frozen object at 0x7e06596ff3d0>,
'decisiontreeregressor__min_samples_leaf':
<scipy.stats._distn_infrastructure.rv_discrete_frozen object at 0x7e065974d990>,
'decisiontreeregressor__min_samples_split':
<scipy.stats._distn_infrastructure.rv_discrete_frozen object at
0x7e065969d720>},
                        random_state=42, scoring='neg_root_mean_squared_error')

[ ]: # Presenting result
tree_uber_rnd_res = pd.DataFrame(tree_uber_rnd_search.cv_results_)
```

```
tree_uber_rnd_res.sort_values(by="mean_test_score", ascending=False,
    inplace=True)
tree_uber_rnd_res.filter(regex = '(^param_|mean_test_score)', axis=1).head(10)
```

```
[ ]: param_decisiontreeregressor__max_depth \
8      12
9      14
5      11
0      16
7      19
4      17
6      19
1      17
2      17
3      17

param_decisiontreeregressor__min_samples_leaf \
8      12
9      19
5      1
0      15
7      16
4      12
6      12
1      7
2      4
3      3

param_decisiontreeregressor__min_samples_split mean_test_score
8      4      -1.933138
9      8      -1.937784
5     13      -1.944010
0     12      -1.959175
7     16      -1.980236
4      7      -1.981772
6     18      -2.001345
1     12      -2.011283
2      9      -2.041862
3      3      -2.070410
```

```
[ ]: # Install necessary libraries
!pip install scikit-optimize
from skopt import BayesSearchCV
from skopt.space import Categorical
```

Requirement already satisfied: scikit-optimize in
/usr/local/lib/python3.10/dist-packages (0.9.0)

Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.10/dist-packages (from scikit-optimize) (1.3.2)
 Requirement already satisfied: pyaml>=16.9 in /usr/local/lib/python3.10/dist-packages (from scikit-optimize) (23.9.7)
 Requirement already satisfied: numpy>=1.13.3 in /usr/local/lib/python3.10/dist-packages (from scikit-optimize) (1.23.5)
 Requirement already satisfied: scipy>=0.19.1 in /usr/local/lib/python3.10/dist-packages (from scikit-optimize) (1.11.4)
 Requirement already satisfied: scikit-learn>=0.20.0 in /usr/local/lib/python3.10/dist-packages (from scikit-optimize) (1.2.2)
 Requirement already satisfied: PyYAML in /usr/local/lib/python3.10/dist-packages (from pyaml>=16.9->scikit-optimize) (6.0.1)
 Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.20.0->scikit-optimize) (3.2.0)

```
[ ]: # DecisionTreeRegressor Bayes Search for Uber dataframe
param_grid = {
    'decisiontreeregressor__max_depth' : [10,20],
    'decisiontreeregressor__min_samples_split' : [2,5,10],
    'decisiontreeregressor__min_samples_leaf' : [1,5,10]
}

tree_uber_bayesian_search = BayesSearchCV(
    estimator=tree_uber,
    search_spaces=param_grid,
    n_iter=10,
    cv=3,
    scoring='neg_root_mean_squared_error',
    random_state=42
)

tree_uber_bayesian_search.fit(X_train_uber, y_train_uber)
```

```
[ ]: BayesSearchCV(cv=3,
                  estimator=Pipeline(steps=[('columntransformer',
                                             ColumnTransformer(transformers=[('cat',
OneHotEncoder(drop='first'),
['cab_type',
'destination',
'source',
'name',
'time_period'])),
                                             ('num',
StandardScaler(),
['distance',
'temp',
```

```

'clouds',
'pressure',
'rain',
'humidity',
'wind']]])),

                                ('decisiontreeregressor',
DecisionTreeRegressor(random_state=42))]],
                                n_iter=10, random_state=42, scoring='neg_root_mean_squared_error',
                                search_spaces={'decisiontreeregressor__max_depth': [10, 20],
                                                'decisiontreeregressor__min_samples_leaf': [1, 5,
                                                                                             10],
                                                'decisiontreeregressor__min_samples_split': [2, 5,
                                                                                             10]})

```

```

[ ]: # Presenting result
tree_uber_bayes_res = pd.DataFrame(tree_uber_bayesian_search.cv_results_)
tree_uber_bayes_res.sort_values(by="mean_test_score", ascending=False,
                                inplace=True)
tree_uber_bayes_res.filter(regex = '(^param_|mean_test_score)', axis=1).head(10)

```

```

[ ]:  param_decisiontreeregressor__max_depth  \
9                                           10
2                                           14
7                                           15
0                                           14
6                                           16
5                                           17
1                                           18
4                                           18
8                                           20
3                                           18

```

```

    param_decisiontreeregressor__min_samples_leaf  \
9                                           10
2                                           10
7                                           10
0                                           5
6                                           10
5                                           10
1                                           10
4                                           5
8                                           5
3                                           1

```

```

    param_decisiontreeregressor__min_samples_split  mean_test_score
9                                           5      -1.935140
2                                           2      -1.953656

```

7	5	-1.965383
0	10	-1.971660
6	5	-1.979480
5	2	-1.992081
1	5	-2.003567
4	5	-2.047339
8	10	-2.082457
3	5	-2.171583

4.2.2 Lyft

```
[ ]: # DecisionTreeRegressor for Lyft dataframe
tree_lyft = make_pipeline(preprocess_pipeline,
    ↳DecisionTreeRegressor(random_state=42))
tree_lyft.fit(X_train_lyft, y_train_lyft)

tree_cv_rmse = -cross_val_score(tree_lyft, X_train_lyft, y_train_lyft,
    scoring="neg_root_mean_squared_error", cv=3)
print(f"Average DecisionTree Regression Cross-Validation RMSE: {tree_cv_rmse.
    ↳mean():.0f}")
```

Average DecisionTree Regression Cross-Validation RMSE: 4

```
[ ]: # DecisionTreeRegressor Grid Search for Lyft dataframe
from sklearn.model_selection import GridSearchCV
import numpy as np

param_grid = {
    'decisiontreeregressor__max_depth' : [10,20],
    'decisiontreeregressor__min_samples_split' : [2,5,10],
    'decisiontreeregressor__min_samples_leaf' : [1,5,10]
}

tree_lyft_grid_search = GridSearchCV(tree_lyft, param_grid, cv=3,
    ↳scoring='neg_root_mean_squared_error')
tree_lyft_grid_search.fit(X_train_lyft, y_train_lyft)
```

```
[ ]: GridSearchCV(cv=3,
    estimator=Pipeline(steps=[('columntransformer',
    ColumnTransformer(transformers=[('cat',
OneHotEncoder(drop='first'),
['cab_type',
'destination',
'source',
'name',
'time_period'])],
    ('num',
```

```

StandardScaler(),
['distance',
 'temp',
 'clouds',
 'pressure',
 'rain',
 'humidity',
 'wind']]])),
                                ('decisiontreeregressor',
DecisionTreeRegressor(random_state=42))]),
                                param_grid={'decisiontreeregressor__max_depth': [10, 20],
                                              'decisiontreeregressor__min_samples_leaf': [1, 5, 10],
                                              'decisiontreeregressor__min_samples_split': [2, 5,
                                                                                             10]},
                                scoring='neg_root_mean_squared_error')

```

```

[ ]: # Presenting result
tree_lyft_cv = pd.DataFrame(tree_lyft_grid_search.cv_results_)
tree_lyft_cv.sort_values(by="mean_test_score", ascending=False, inplace=True)
tree_lyft_cv.filter(regex = '(^param_|mean_test_score)', axis=1).head(10)

```

```

[ ]:  param_decisiontreeregressor__max_depth  \
6                                           10
7                                           10
8                                           10
3                                           10
4                                           10
5                                           10
2                                           10
1                                           10
0                                           10
15                                          20

      param_decisiontreeregressor__min_samples_leaf  \
6                                                    10
7                                                    10
8                                                    10
3                                                    5
4                                                    5
5                                                    5
2                                                    1
1                                                    1
0                                                    1
15                                                  10

      param_decisiontreeregressor__min_samples_split  mean_test_score
6                                                    2          -3.109589

```

7	5	-3.109589
8	10	-3.109589
3	2	-3.121271
4	5	-3.121271
5	10	-3.121271
2	10	-3.143029
1	5	-3.152581
0	2	-3.162694
15	2	-3.275394

```
[ ]: # DecisionTreeRegressor Random Search for Lyft dataframe
```

```
from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import randint

param_grid = {
    'decisiontreeregressor__max_depth' : randint(10,20),
    'decisiontreeregressor__min_samples_split' : randint(2,20),
    'decisiontreeregressor__min_samples_leaf' : randint(1,20)
}

tree_lyft_rnd_search = RandomizedSearchCV(
    tree_lyft, param_grid, n_iter=10, cv=3,
    scoring='neg_root_mean_squared_error', random_state=42)

tree_lyft_rnd_search.fit(X_train_lyft, y_train_lyft)
```

```
[ ]: RandomizedSearchCV(cv=3,
                        estimator=Pipeline(steps=[('columntransformer',
ColumnTransformer(transformers=[('cat',
OneHotEncoder(drop='first'),
['cab_type',
'destination',
'source',
'name',
'time_period'])),
('num',
StandardScaler()),
['distance',
'temp',
'clouds',
'pressure',
'rain',
'humidity',
'wind'])])),
                        ('decisiontreeregressor',
DecisionTreeRegressor(random_state=4...
                        param_distributions={'decisiontreeregressor__max_depth':
```

```
<scipy.stats._distn_infrastructure.rv_discrete_frozen object at 0x7e06596fcf0>,
'decisiontreeregressor__min_samples_leaf':
<scipy.stats._distn_infrastructure.rv_discrete_frozen object at 0x7e0659768850>,
'decisiontreeregressor__min_samples_split':
<scipy.stats._distn_infrastructure.rv_discrete_frozen object at
0x7e065974f160>},
    random_state=42, scoring='neg_root_mean_squared_error')
```

```
[ ]: # Presenting result
tree_lyft_rnd_res = pd.DataFrame(tree_lyft_rnd_search.cv_results_)
tree_lyft_rnd_res.sort_values(by="mean_test_score", ascending=False,
    inplace=True)
tree_lyft_rnd_res.filter(regex = '(^param_|mean_test_score)', axis=1).head(10)
```

```
[ ]: param_decisiontreeregressor__max_depth \
9      14
8      12
5      11
0      16
7      19
4      17
6      19
1      17
2      17
3      17

param_decisiontreeregressor__min_samples_leaf \
9      19
8      12
5      1
0      15
7      16
4      12
6      12
1      7
2      4
3      3

param_decisiontreeregressor__min_samples_split mean_test_score
9      8      -3.123658
8      4      -3.129115
5     13      -3.155909
0     12      -3.175492
7     16      -3.200898
4      7      -3.218256
6     18      -3.241673
1     12      -3.285188
```

2	9	-3.374799
3	3	-3.442527

```
[ ]: # Install necessary libraries
!pip install scikit-optimize
from skopt import BayesSearchCV
from skopt.space import Categorical
```

```
Requirement already satisfied: scikit-optimize in
/usr/local/lib/python3.10/dist-packages (0.9.0)
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.10/dist-
packages (from scikit-optimize) (1.3.2)
Requirement already satisfied: pyaml>=16.9 in /usr/local/lib/python3.10/dist-
packages (from scikit-optimize) (23.9.7)
Requirement already satisfied: numpy>=1.13.3 in /usr/local/lib/python3.10/dist-
packages (from scikit-optimize) (1.23.5)
Requirement already satisfied: scipy>=0.19.1 in /usr/local/lib/python3.10/dist-
packages (from scikit-optimize) (1.11.4)
Requirement already satisfied: scikit-learn>=0.20.0 in
/usr/local/lib/python3.10/dist-packages (from scikit-optimize) (1.2.2)
Requirement already satisfied: PyYAML in /usr/local/lib/python3.10/dist-packages
(from pyaml>=16.9->scikit-optimize) (6.0.1)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.20.0->scikit-
optimize) (3.2.0)
```

```
[ ]: # DecisionTreeRegressor Bayes Search for Lyft dataframe
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import cross_val_score
from sklearn.pipeline import Pipeline
from scipy.stats import randint

param_grid = {
    'decisiontreeregressor__max_depth' : [10,20],
    'decisiontreeregressor__min_samples_split' : [2,5,10],
    'decisiontreeregressor__min_samples_leaf' : [1,5,10]
}

tree_lyft_bayesian_search = BayesSearchCV(
    estimator=tree_lyft,
    search_spaces=param_grid,
    n_iter=10,
    cv=3,
    scoring='neg_root_mean_squared_error',
    random_state=42
)
```

```
tree_lyft_bayesian_search.fit(X_train_lyft, y_train_lyft)
```

```
[ ]: BayesSearchCV(cv=3,
                  estimator=Pipeline(steps=[('columntransformer',
                                             ColumnTransformer(transformers=[('cat',
OneHotEncoder(drop='first'),
['cab_type',
'destination',
'source',
'name',
'time_period'])),
                                             ('num',
StandardScaler(),
['distance',
'temp',
'clouds',
'pressure',
'rain',
'humidity',
'wind'])])),
                  DecisionTreeRegressor(random_state=42)),
                  n_iter=10, random_state=42, scoring='neg_root_mean_squared_error',
                  search_spaces={'decisiontreeregressor__max_depth': [10, 20],
                                'decisiontreeregressor__min_samples_leaf': [1, 5,
                                                                              10],
                                'decisiontreeregressor__min_samples_split': [2, 5,
                                                                              10]})
```

```
[ ]: # Presenting result
tree_lyft_bayes_res = pd.DataFrame(tree_lyft_bayesian_search.cv_results_)
tree_lyft_bayes_res.sort_values(by="mean_test_score", ascending=False,
                                ↪inplace=True)
tree_lyft_bayes_res.filter(regex = '(^param_|mean_test_score)', axis=1).head(10)
```

```
[ ]: param_decisiontreeregressor__max_depth \
9      10
2      14
7      15
6      16
5      17
0      14
1      18
4      18
8      20
3      18
```


	param_decisiontreeregressor__min_samples_leaf \
9	10
2	10
7	10
6	10
5	10
0	5
1	10
4	5
8	5
3	1

	param_decisiontreeregressor__min_samples_split	mean_test_score
9	5	-3.109589
2	2	-3.174279
7	5	-3.196858
6	5	-3.217336
5	2	-3.235908
0	10	-3.238037
1	5	-3.250482
4	5	-3.377004
8	10	-3.425158
3	5	-3.675220

4.3 Elastic Net Regression

4.3.1 Uber

```
[ ]: # Elastic Net regression for Uber dataframe
from sklearn.linear_model import ElasticNet
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import cross_val_score

elastic_net_uber = make_pipeline(preprocess_pipeline, ElasticNet())
elastic_net_uber.fit(X_train_uber, y_train_uber)
elastic_net_cv_rmse = -cross_val_score(elastic_net_uber, X_train_uber,
    ↪ y_train_uber,
                                   scoring="neg_root_mean_squared_error", cv=3)
print(f"Average Elastic Net Regression Cross-Validation RMSE:
    ↪ {elastic_net_cv_rmse.mean():.0f}")
```

Average Elastic Net Regression Cross-Validation RMSE: 7

```
[ ]: # Elastic Net Regression Grid Search for Uber dataframe
from sklearn.model_selection import GridSearchCV
```

```

from sklearn.pipeline import Pipeline

param_grid = {
    'elasticnet__alpha': [0.1, 0.5, 1.0],
    'elasticnet__l1_ratio': [0.1, 0.5, 0.9],
    'elasticnet__fit_intercept': [True, False],
    'elasticnet__positive': [True, False],}

elastic_net_uber = make_pipeline(preprocess_pipeline, ElasticNet())
elastic_net_grid_search = GridSearchCV(elastic_net_uber, param_grid, cv=3,
    ↪scoring='neg_root_mean_squared_error')
elastic_net_grid_search.fit(X_train_uber, y_train_uber)

```

```

[ ]: GridSearchCV(cv=3,
                estimator=Pipeline(steps=[('columntransformer',
                                           ColumnTransformer(transformers=[('cat',
OneHotEncoder(drop='first'),
['cab_type',
'destination',
'source',
'name',
'time_period'])),
                                           ('num',
StandardScaler(),
['distance',
'temp',
'clouds',
'pressure',
'rain',
'humidity',
'wind'])])),
                ('elasticnet', ElasticNet()))],
        param_grid={'elasticnet__alpha': [0.1, 0.5, 1.0],
                    'elasticnet__fit_intercept': [True, False],
                    'elasticnet__l1_ratio': [0.1, 0.5, 0.9],
                    'elasticnet__positive': [True, False]},
        scoring='neg_root_mean_squared_error')

```

```

[ ]: # Presenting result
elastic_net_uber_cv = pd.DataFrame(elastic_net_grid_search.cv_results_)
elastic_net_uber_cv.sort_values(by="mean_test_score", ascending=False,
    ↪inplace=True)
elastic_net_uber_cv.filter(regex='(^param_|mean_test_score)', axis=1).head(10)

```

```

[ ]:   param_elasticnet__alpha  param_elasticnet__fit_intercept \
5                0.1                True
3                0.1                True

```

1	0.1	True
17	0.5	True
4	0.1	True
2	0.1	True
11	0.1	False
10	0.1	False
0	0.1	True
16	0.5	True

	param_elasticnet__l1_ratio	param_elasticnet__positive	mean_test_score
5	0.9	False	-2.773754
3	0.5	False	-3.413798
1	0.1	False	-3.928199
17	0.9	False	-4.622299
4	0.9	True	-4.695244
2	0.5	True	-5.032953
11	0.9	False	-5.065001
10	0.9	True	-5.282374
0	0.1	True	-5.374240
16	0.9	True	-5.484940

```
[ ]: # Elastic Net Regression Random Search for Uber dataframe
from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import uniform

param_distributions = {
    'elasticnet__alpha': uniform(0, 1),
    'elasticnet__l1_ratio': uniform(0, 1),
    'elasticnet__fit_intercept': [True, False],
    'elasticnet__positive': [True, False]}

elastic_net_rnd_search = RandomizedSearchCV(
    elastic_net_uber, param_distributions=param_distributions, n_iter=10, cv=3,
    scoring='neg_root_mean_squared_error', random_state=42)

elastic_net_rnd_search.fit(X_train_uber, y_train_uber)
```

```
[ ]: RandomizedSearchCV(cv=3,
                        estimator=Pipeline(steps=[('columntransformer',
ColumnTransformer(transformers=[('cat',
OneHotEncoder(drop='first'),
['cab_type',
'destination',
'source',
'name',
'time_period'])],
('num',
```

```

StandardScaler(),
['distance',
 'temp',
 'clouds',
 'pressure',
 'rain',
 'humidity',
 'wind']]])),
                                ('elasticnet', ElasticNet()))],
                                param_distributions={'elasticnet__alpha':
<scipy.stats._distn_infrastructure.rv_continuous_frozen object at
0x7e06595be920>,
                                'elasticnet__fit_intercept': [True,
                                                                False],
                                'elasticnet__l1_ratio':
<scipy.stats._distn_infrastructure.rv_continuous_frozen object at
0x7e065976acb0>,
                                'elasticnet__positive': [True, False]},
                                random_state=42, scoring='neg_root_mean_squared_error')

```

```

[ ]: # Presenting result
elastic_net_uber_rnd_res = pd.DataFrame(elastic_net_rnd_search.cv_results_)
elastic_net_uber_rnd_res.sort_values(by="mean_test_score", ascending=False,
    ↪inplace=True)
elastic_net_uber_rnd_res.filter(regex='^(param_|mean_test_score)').head(10)

```

```

[ ]:  param_elasticnet__alpha  param_elasticnet__fit_intercept  \
2                0.058084                                False
7                0.139494                                False
8                0.45607                                 True
0                0.37454                                 True
6                0.431945                                True
5                0.183405                                False
1                0.598658                                True
9                0.514234                                False
3                0.708073                                False
4                0.832443                                False

    param_elasticnet__l1_ratio  param_elasticnet__positive  mean_test_score
2                0.333709                                False        -5.316421
7                0.973756                                 True        -5.386674
8                0.618386                                False        -5.438361
0                0.183435                                False        -5.849585
6                0.524775                                 True        -6.352243
5                0.611653                                 True        -6.481434
1                0.445833                                 True        -6.832557
9                0.466763                                 True       -10.021694

```

3	0.056412	False	-12.272700
4	0.000779	False	-12.892390

```
[ ]: # Elastic Net Regression Bayes Search for Uber dataframe
sparse_columns = ['cab_type', 'destination', 'source', 'name', 'time_period']

X_train_uber_selected = X_train_uber[sparse_columns].to_numpy()
uber_bayes_encoder = OneHotEncoder(sparse=False)
X_train_uber_encoded = pd.get_dummies(X_train_uber[sparse_columns])
X_train_uber_bayes_final = pd.concat([X_train_uber.
    ↪drop(columns=sparse_columns), X_train_uber_encoded], axis=1)

param_distribs = {
    'alpha': (0.01, 1.0, 'uniform'),
    'l1_ratio': (0.01, 1.0, 'uniform'),
    'fit_intercept': [True, False],
    'positive': [True, False]}

elastic_net_bayes_search = BayesSearchCV(
    estimator=ElasticNet(),
    search_spaces=param_distribs,
    n_iter=10, cv=3,
    scoring='neg_root_mean_squared_error', random_state=42)

elastic_net_bayes_search.fit(X_train_uber_bayes_final, y_train_uber)
```

```
[ ]: BayesSearchCV(cv=3, estimator=ElasticNet(), n_iter=10, random_state=42,
    scoring='neg_root_mean_squared_error',
    search_spaces={'alpha': (0.01, 1.0, 'uniform'),
        'fit_intercept': [True, False],
        'l1_ratio': (0.01, 1.0, 'uniform'),
        'positive': [True, False]})
```

```
[ ]: # Presenting result
elastic_net_uber_bayes_search = pd.DataFrame(elastic_net_bayes_search.
    ↪cv_results_)
elastic_net_uber_bayes_search.sort_values(by="mean_test_score",
    ↪ascending=False, inplace=True)
elastic_net_uber_bayes_search.filter(regex='^(param_|mean_test_score)').head(10)
```

```
[ ]: param_alpha param_fit_intercept param_l1_ratio param_positive \
9    0.013594                True      0.74387        False
0    0.416003                True      0.933539       False
2    0.450384                True      0.113811       False
8    0.955923                True      0.872935       False
7    0.547969                True      0.501502        True
5    0.736688                True      0.171971       False
```

6	0.620909	True	0.365433	True
3	0.814272	False	0.602067	True
4	0.801558	False	0.531354	True
1	0.839014	True	0.310376	True

	mean_test_score
9	-2.417090
0	-3.780388
2	-6.048618
8	-6.294131
7	-6.410722
5	-6.702840
6	-6.768686
3	-6.811925
4	-6.893178
1	-7.158599

4.3.2 Lyft

```
[ ]: # Elastic Net regression for Lyft dataframe

elastic_net_lyft = make_pipeline(preprocess_pipeline, ElasticNet())
elastic_net_lyft.fit(X_train_lyft, y_train_lyft)
y_train_predictions_lyft = elastic_net_lyft.predict(X_train_lyft)
elastic_net_cv_rmse = -cross_val_score(elastic_net_lyft, X_train_lyft,
    ↪ y_train_lyft,
                                   scoring="neg_root_mean_squared_error", cv=3)
print(f"Average Elastic Net Regression Cross-Validation RMSE:␣
    ↪ {elastic_net_cv_rmse.mean():.0f}")
```

Average Elastic Net Regression Cross-Validation RMSE: 8

```
[ ]: # Elastic Net Regression Grid Search for Lyft dataframe

param_grid = {
    'elasticnet__alpha': [0.1, 0.5, 1.0],
    'elasticnet__l1_ratio': [0.1, 0.5, 0.9],
    'elasticnet__fit_intercept': [True, False],
    'elasticnet__positive': [True, False],}

elastic_net_lyft = make_pipeline(preprocess_pipeline, ElasticNet())
elastic_net_grid_search = GridSearchCV(elastic_net_lyft, param_grid, cv=3,
    ↪ scoring='neg_root_mean_squared_error')
elastic_net_grid_search.fit(X_train_lyft, y_train_lyft)
```

```
[ ]: GridSearchCV(cv=3,
                  estimator=Pipeline(steps=[('columntransformer',
                                             ColumnTransformer(transformers=[('cat',
OneHotEncoder(drop='first'),
['cab_type',
'destination',
'source',
'name',
'time_period'])),
                                             ('num',
StandardScaler(),
['distance',
'temp',
'clouds',
'pressure',
'rain',
'humidity',
'wind'])])),
                  ('elasticnet', ElasticNet()))],
        param_grid={'elasticnet__alpha': [0.1, 0.5, 1.0],
                    'elasticnet__fit_intercept': [True, False],
                    'elasticnet__l1_ratio': [0.1, 0.5, 0.9],
                    'elasticnet__positive': [True, False]},
        scoring='neg_root_mean_squared_error')
```

```
[ ]: # Presenting result
elastic_net_lyft_cv = pd.DataFrame(elastic_net_grid_search.cv_results_)
elastic_net_lyft_cv.sort_values(by="mean_test_score", ascending=False,
                               inplace=True)
elastic_net_lyft_cv.filter(regex='(^param_|mean_test_score)', axis=1).head(10)
```

```
[ ]:   param_elasticnet__alpha param_elasticnet__fit_intercept \
5          0.1          True
3          0.1          True
1          0.1          True
4          0.1          True
11         0.1         False
17         0.5          True
10         0.1         False
2          0.1          True
9          0.1         False
8          0.1         False

   param_elasticnet__l1_ratio param_elasticnet__positive  mean_test_score
5              0.9          False      -3.643614
3              0.5          False      -4.146584
1              0.1          False      -4.663989
```

4	0.9	True	-5.081580
11	0.9	False	-5.118657
17	0.9	False	-5.144882
10	0.9	True	-5.497436
2	0.5	True	-5.663903
9	0.5	False	-5.842750
8	0.5	True	-6.127570

```
[ ]: # Elastic Net Regression Random Search for Lyft dataframe
```

```
param_distributions = {
    'elasticnet__alpha': uniform(0, 1),
    'elasticnet__l1_ratio': uniform(0, 1),
    'elasticnet__fit_intercept': [True, False],
    'elasticnet__positive': [True, False]}

elastic_net_rnd_search = RandomizedSearchCV(
    elastic_net_lyft, param_distributions=param_distributions, n_iter=10, cv=3,
    scoring='neg_root_mean_squared_error', random_state=42)

elastic_net_rnd_search.fit(X_train_lyft, y_train_lyft)
```

```
[ ]: RandomizedSearchCV(cv=3,
                        estimator=Pipeline(steps=[('columntransformer',
ColumnTransformer(transformers=[('cat',
OneHotEncoder(drop='first'),
['cab_type',
'destination',
'source',
'name',
'time_period'])),
('num',
StandardScaler(),
['distance',
'temp',
'clouds',
'pressure',
'rain',
'humidity',
'wind'])])),
                        ('elasticnet', ElasticNet()))],
                        param_distributions={'elasticnet__alpha':
<scipy.stats._distn_infrastructure.rv_continuous_frozen object at
0x7e06595bc7c0>,
                        'elasticnet__fit_intercept': [True,
False],
                        'elasticnet__l1_ratio':
```



```
<scipy.stats._distn_infrastructure.rv_continuous_frozen object at
0x7e0659612110>,
                                'elasticnet__positive': [True, False]},
                                random_state=42, scoring='neg_root_mean_squared_error')
```

```
[ ]: # Presenting result
elastic_net_lyft_rnd_res = pd.DataFrame(elastic_net_rnd_search.cv_results_)
elastic_net_lyft_rnd_res.sort_values(by="mean_test_score", ascending=False,
    ↪ inplace=True)
elastic_net_lyft_rnd_res.filter(regex='^(param_|mean_test_score)').head(10)
```

```
[ ]: param_elasticnet__alpha param_elasticnet__fit_intercept \
2                0.058084                False
7                0.139494                False
8                0.45607                 True
0                0.37454                 True
5                0.183405                False
6                0.431945                 True
1                0.598658                 True
9                0.514234                False
3                0.708073                False
4                0.832443                False

param_elasticnet__l1_ratio param_elasticnet__positive mean_test_score
2                0.333709                False        -5.401231
7                0.973756                 True        -5.634315
8                0.618386                False        -6.244404
0                0.183435                False        -6.774094
5                0.611653                 True        -6.793337
6                0.524775                 True        -7.433733
1                0.445833                 True        -8.012347
9                0.466763                 True       -10.750987
3                0.056412                False       -13.474418
4                0.000779                False       -14.204772
```

```
[ ]: # Elastic Net Regression Bayes Search for Lyft dataframe
X_train_lyft_selected = X_train_lyft[sparse_columns].to_numpy()
lyft_bayes_encoder = OneHotEncoder(sparse=False)
X_train_lyft_encoded = pd.get_dummies(X_train_lyft[sparse_columns])
X_train_lyft_bayes_final = pd.concat([X_train_lyft.
    ↪ drop(columns=sparse_columns), X_train_lyft_encoded], axis=1)

param_distribs = {
    'alpha': (0.01, 1.0, 'uniform'),
    'l1_ratio': (0.01, 1.0, 'uniform'),
    'fit_intercept': [True, False],
    'positive': [True, False]}
```

```

elastic_net_bayes_search = BayesSearchCV(
    estimator=ElasticNet(),
    search_spaces=param_distributions,
    n_iter=10, cv=3,
    scoring='neg_root_mean_squared_error', random_state=42)

elastic_net_bayes_search.fit(X_train_lyft_bayes_final, y_train_lyft)

```

```

[ ]: BayesSearchCV(cv=3, estimator=ElasticNet(), n_iter=10, random_state=42,
    scoring='neg_root_mean_squared_error',
    search_spaces={'alpha': (0.01, 1.0, 'uniform'),
        'fit_intercept': [True, False],
        'l1_ratio': (0.01, 1.0, 'uniform'),
        'positive': [True, False]})

```

```

[ ]: # Presenting result
elastic_net_lyft_bayes_search = pd.DataFrame(elastic_net_bayes_search.
    cv_results_)
elastic_net_lyft_bayes_search.sort_values(by="mean_test_score",
    ascending=False, inplace=True)
elastic_net_lyft_bayes_search.filter(regex='^(param_|mean_test_score)').head(10)

```

```

[ ]:  param_alpha  param_fit_intercept  param_l1_ratio  param_positive  \
9      0.013594                True      0.74387        False
0      0.416003                True      0.933539        False
8      0.955923                True      0.872935        False
2      0.450384                True      0.113811        False
7      0.547969                True      0.501502         True
5      0.736688                True      0.171971        False
6      0.620909                True      0.365433         True
3      0.814272               False      0.602067         True
4      0.801558               False      0.531354         True
1      0.839014                True      0.310376         True

    mean_test_score
9      -3.491303
0      -4.584596
8      -7.000939
2      -7.170831
7      -7.796131
5      -7.880016
6      -8.139924
3      -8.169614
4      -8.252276
1      -8.520885

```

4.4 XGBoost Regression

4.4.1 Uber

```
[ ]: pip install xgboost
```

```
Requirement already satisfied: xgboost in /usr/local/lib/python3.10/dist-packages (2.0.2)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from xgboost) (1.23.5)
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (from xgboost) (1.11.4)
```

```
[ ]: # XGBoost Regression for Uber dataframe
import xgboost as xgb
from xgboost import XGBRegressor
from sklearn.metrics import accuracy_score, mean_squared_error
import numpy as np

xgb_uber = make_pipeline(preprocess_pipeline, XGBRegressor())
xgb_uber.fit(X_train_uber, y_train_uber)
y_pred_uber = xgb_uber.predict(X_test_uber)
xgb_rmse = -cross_val_score(xgb_uber, X_train_uber, y_train_uber,
                           scoring="neg_root_mean_squared_error", cv=3)
print(f"Average XGB Cross-Validation RMSE: {xgb_rmse.mean():.0f}")
```

Average XGB Cross-Validation RMSE: 2

```
[ ]: # XGB Grid Search for Uber dataframe
param_grid = {
    'xgbregressor__n_estimators': [50, 100],
    'xgbregressor__max_depth': [3, 5],
    'xgbregressor__colsample_bytree': [0.8, 1.0]
}

xgb_uber_grid_search = GridSearchCV(xgb_uber, param_grid, cv=2,
    scoring='neg_root_mean_squared_error', n_jobs=-1)
xgb_uber_grid_search.fit(X_train_uber, y_train_uber)

#Evaluate Performance
y_pred_grid = xgb_uber_grid_search.predict(X_test_uber)
rmse_grid = np.sqrt(mean_squared_error(y_test_uber, y_pred_grid))
print(f"Grid Search RMSE: {rmse_grid}")
```

Grid Search RMSE: 1.8760529844851286

```
[ ]: # Presenting result
xgb_uber_grid = pd.DataFrame(xgb_uber_grid_search.cv_results_)
xgb_uber_grid.sort_values(by="mean_test_score", ascending=False, inplace=True)
xgb_uber_grid.filter(regex='(^param_|mean_test_score)', axis=1).head(10)
```

```
[ ]:  param_xgbregressor__colsample_bytree param_xgbregressor__max_depth \
3          0.8          5
7          1.0          5
2          0.8          5
6          1.0          5
5          1.0          3
1          0.8          3
4          1.0          3
0          0.8          3

    param_xgbregressor__n_estimators  mean_test_score
3                100      -1.889813
7                100      -1.892140
2                 50      -1.906587
6                 50      -1.906973
5                100      -1.942402
1                100      -1.948609
4                 50      -2.026792
0                 50      -2.029297
```

```
[ ]: # XGB Random Search for Uber dataframe
param_grid = {
    'xgbregressor__n_estimators': [50, 100],
    'xgbregressor__max_depth': [3, 5],
    'xgbregressor__colsample_bytree': [0.8, 1.0]
}

xgb_uber_Random_search = RandomizedSearchCV(xgb_uber, param_grid, cv=2,
    ↪scoring='neg_root_mean_squared_error', n_jobs=-1)
xgb_uber_Random_search.fit(X_train_uber, y_train_uber)

#Evaluate Performance
y_pred_Random = xgb_uber_Random_search.predict(X_test_uber)
rmse_Random = np.sqrt(mean_squared_error(y_test_uber, y_pred_Random))
print(f"Random Search RMSE: {rmse_Random}")
```

Random Search RMSE: 1.8760529844851286

```
[ ]: # Presenting result
xgb_uber_Random = pd.DataFrame(xgb_uber_Random_search.cv_results_)
xgb_uber_Random.sort_values(by="mean_test_score", ascending=False, inplace=True)
xgb_uber_Random.filter(regex='(^param_|mean_test_score)', axis=1).head(10)
```

```
[ ]: param_xgbregressor__n_estimators param_xgbregressor__max_depth \
3                                100                                5
7                                100                                5
2                                50                                 5
6                                50                                 5
5                                100                                3
1                                100                                3
4                                50                                 3
0                                50                                 3

param_xgbregressor__colsample_bytree mean_test_score
3                                0.8            -1.889813
7                                1.0            -1.892140
2                                0.8            -1.906587
6                                1.0            -1.906973
5                                1.0            -1.942402
1                                0.8            -1.948609
4                                1.0            -2.026792
0                                0.8            -2.029297
```

```
[ ]: # XGB Bayes Search for Uber dataframe
param_bayes = {
    'xgbregressor__n_estimators': (50, 200),
    'xgbregressor__max_depth': (3, 7),
    'xgbregressor__colsample_bytree': (0.8, 1.0)
}

xgb_uber_bayes_search = BayesSearchCV(xgb_uber, search_spaces=param_bayes,
    n_iter=10, cv=2, scoring='neg_root_mean_squared_error', n_jobs=-1)
xgb_uber_bayes_search.fit(X_train_uber, y_train_uber)

# Evaluate Performance - Bayes Search
y_pred_bayes = xgb_uber_bayes_search.predict(X_test_uber)
rmse_bayes = np.sqrt(mean_squared_error(y_test_uber, y_pred_bayes))
print(f"Bayes Search RMSE: {rmse_bayes}")
```

Bayes Search RMSE: 1.8650364479779769

```
[ ]: # Presenting result
xgb_uber_bayes = pd.DataFrame(xgb_uber_bayes_search.cv_results_)
xgb_uber_bayes.sort_values(by="mean_test_score", ascending=False, inplace=True)
xgb_uber_bayes.filter(regex='(^param_|mean_test_score)', axis=1).head(10)
```

```
[ ]: param_xgbregressor__colsample_bytree param_xgbregressor__max_depth \
3                                0.900701                                5
7                                0.892691                                6
9                                0.809514                                6
```

5	0.931857	7
4	0.935327	4
0	0.860193	7
6	0.857183	7
1	0.903436	7
8	0.903383	4
2	0.885276	4

	param_xgbregressor__n_estimators	mean_test_score
3	151	-1.886856
7	145	-1.887563
9	106	-1.889392
5	84	-1.891292
4	147	-1.895702
0	125	-1.896715
6	154	-1.899093
1	191	-1.902590
8	90	-1.905684
2	54	-1.930349

4.4.2 Lyft

```
[ ]: # XGBoost Regression for Lyft dataframe
xgb_lyft = make_pipeline(preprocess_pipeline, XGBRegressor())
xgb_lyft.fit(X_train_lyft, y_train_lyft)
y_pred_lyft = xgb_lyft.predict(X_test_lyft)
xgb_rmse = -cross_val_score(xgb_lyft, X_train_lyft, y_train_lyft,
                             scoring="neg_root_mean_squared_error", cv=3)
print(f"Average XGB Cross-Validation RMSE: {xgb_rmse.mean():.0f}")
```

Average XGB Cross-Validation RMSE: 3

```
[ ]: # XGB Grid Search for lyft dataframe
param_grid = {
    'xgbregressor__n_estimators': [50, 100],
    'xgbregressor__max_depth': [3, 5],
    'xgbregressor__colsample_bytree': [0.8, 1.0]
}

xgb_lyft_grid_search = GridSearchCV(xgb_lyft, param_grid, cv=2,
                                     scoring='neg_root_mean_squared_error', n_jobs=-1)
xgb_lyft_grid_search.fit(X_train_lyft, y_train_lyft)

#Evaluate Performance
y_pred_grid = xgb_lyft_grid_search.predict(X_test_lyft)
rmse_grid = np.sqrt(mean_squared_error(y_test_lyft, y_pred_grid))
print(f"Grid Search RMSE: {rmse_grid}")
```

Grid Search RMSE: 2.9914378947719418

```
[ ]: # Presenting result
xgb_lyft_grid = pd.DataFrame(xgb_lyft_grid_search.cv_results_)
xgb_lyft_grid.sort_values(by="mean_test_score", ascending=False, inplace=True)
xgb_lyft_grid.filter(regex='(^param_|mean_test_score)', axis=1).head(10)
```

```
[ ]:   param_xgbregressor__colsample_bytree param_xgbregressor__max_depth \
3                                0.8                                5
7                                1.0                                5
6                                1.0                                5
2                                0.8                                5
5                                1.0                                3
1                                0.8                                3
4                                1.0                                3
0                                0.8                                3

   param_xgbregressor__n_estimators mean_test_score
3                                100      -2.990797
7                                100      -2.991257
6                                 50      -3.022386
2                                 50      -3.024419
5                                100      -3.052865
1                                100      -3.064362
4                                 50      -3.097627
0                                 50      -3.113629
```

```
[ ]: # XGB Random Search for lyft dataframe
param_grid = {
    'xgbregressor__n_estimators': [50, 100],
    'xgbregressor__max_depth': [3, 5],
    'xgbregressor__subsample': [0.8, 1.0],
}

xgb_lyft_Random_search = RandomizedSearchCV(xgb_lyft, param_grid, cv=2,
                                             scoring='neg_root_mean_squared_error', n_jobs=-1)
xgb_lyft_Random_search.fit(X_train_lyft, y_train_lyft)

#Evaluate Performance
y_pred_Random = xgb_lyft_Random_search.predict(X_test_lyft)
rmse_Random = np.sqrt(mean_squared_error(y_test_lyft, y_pred_Random))
print(f"Random Search RMSE: {rmse_Random}")
```

Random Search RMSE: 2.9910621413885305

```
[ ]: # Presenting result
xgb_lyft_Random = pd.DataFrame(xgb_lyft_Random_search.cv_results_)
```

```
xgb_lyft_Random.sort_values(by="mean_test_score", ascending=False, inplace=True)
xgb_lyft_Random.filter(regex='(^param_|mean_test_score)', axis=1).head(10)
```

```
[ ]:  param_xgbregressor__subsample param_xgbregressor__n_estimators \
7          1.0                      100
6          0.8                      100
4          0.8                       50
5          1.0                       50
2          0.8                      100
3          1.0                      100
0          0.8                       50
1          1.0                       50

param_xgbregressor__max_depth mean_test_score
7          5          -2.991257
6          5          -2.998062
4          5          -3.019814
5          5          -3.022386
2          3          -3.051192
3          3          -3.052865
0          3          -3.093635
1          3          -3.097627
```

```
[ ]: # XGB Bayes Search for lyft dataframe
param_bayes = {
    'xgbregressor__n_estimators': (50, 200),
    'xgbregressor__max_depth': (3, 7),
    'xgbregressor__subsample': (0.8, 1.0),
}

xgb_lyft_bayes_search = BayesSearchCV(xgb_lyft, search_spaces=param_bayes,
    n_iter=10, cv=2, scoring='neg_root_mean_squared_error', n_jobs=-1)
xgb_lyft_bayes_search.fit(X_train_lyft, y_train_lyft)

# Evaluate Performance - Bayes Search
y_pred_bayes = xgb_lyft_bayes_search.predict(X_test_lyft)
rmse_bayes = np.sqrt(mean_squared_error(y_test_lyft, y_pred_bayes))
print(f"Bayes Search RMSE: {rmse_bayes}")
```

Bayes Search RMSE: 2.945334385933961

```
[ ]: # Presenting result
xgb_lyft_bayes = pd.DataFrame(xgb_lyft_bayes_search.cv_results_)
xgb_lyft_bayes.sort_values(by="mean_test_score", ascending=False, inplace=True)
xgb_lyft_bayes.filter(regex='(^param_|mean_test_score)', axis=1).head(10)
```



```
[ ]: param_xgbregressor__max_depth param_xgbregressor__n_estimators \
9          5          198
8          6          82
3          6          79
6          6          62
2          5          113
1          4          126
7          5          74
5          5          55
0          4          66
4          3          103

param_xgbregressor__subsample mean_test_score
9          0.946131      -2.964890
8          0.977843      -2.980644
3          0.986264      -2.982536
6          0.891606      -2.986291
2          0.880322      -2.987961
1          0.810927      -3.006045
7          0.878889      -3.006896
5          0.86137       -3.012315
0          0.823544      -3.034166
4          0.933904      -3.053316
```

4.5 Voting Classifier

```
[ ]: from sklearn.ensemble import VotingRegressor
from sklearn.metrics import mean_squared_error
import numpy as np

# Create Uber Voting Classifier
voting_regressor_uber = VotingRegressor(estimators=[
    ('logreg', lin_reg_uber),
    ('decision_tree', tree_uber),
    ('xgboost', xgb_uber),
    ('elastic_net', elastic_net_uber)
])

# Fit and predict the voting classifier
voting_regressor_uber.fit(X_train_uber, y_train_uber)
uber_predictions = voting_regressor_uber.predict(X_test_uber)

# Calculate RMSE
rmse = -np.sqrt(mean_squared_error(y_test_uber, uber_predictions))
```

```

print(f'Voting Regressor RMSE: {rmse}')

# Create Lyft Voting Classifier
voting_regressor_lyft = VotingRegressor(estimators=[
    ('logreg', lin_reg_lyft),
    ('decision_tree', tree_lyft),
    ('xgboost', xgb_lyft),
    ('elastic_net', elastic_net_lyft)
])

# Fit and predict the voting classifier
voting_regressor_lyft.fit(X_train_lyft, y_train_lyft)
lyft_predictions = voting_regressor_lyft.predict(X_test_lyft)

# Calculate RMSE
rmse = -np.sqrt(mean_squared_error(y_test_lyft, lyft_predictions))
print(f'Voting Regressor RMSE: {rmse}')

```

Voting Regressor RMSE: -2.6287571923463147

Voting Regressor RMSE: -3.6599796659916417

Since voting classifier takes the average RMSE of the four models, we cannot identify the best model. We decided to go forward with hyper-parameter tuning.

5 Feature importance

To determine the best model for price prediction for Uber and Lyft, we will select the model with the maximum Negative Root Mean Squared Error (NRMSE).

```

[ ]: import pandas as pd

dataframes_uber = [lin_uber_cv,
                    lin_uber_rnd_res,
                    lin_uber_bayes_res,
                    tree_uber_cv,
                    tree_uber_rnd_res,
                    tree_uber_bayes_res,
                    elastic_net_uber_cv,
                    elastic_net_uber_rnd_res,
                    elastic_net_uber_bayes_search,
                    xgb_uber_grid,
                    xgb_uber_Random,
                    xgb_uber_bayes]

# Loop through the 12 models and find the maximum mean_test_score
for df in dataframes_uber:
    df.reset_index(drop=True,inplace=True)

```

```

subset_values_uber = [df.loc[0, 'mean_test_score'] for df in dataframes_uber]

max_value_uber = max(subset_values_uber)
dataframe_name = [name for name, obj in globals().items() if obj is_
↳dataframes_uber[subset_values_uber.index(max_value_uber)]] [0]

print("Best Model for Uber prediction is", dataframe_name, ": ", max_value_uber)

dataframes_lyft = [lin_lyft_cv,
                    lin_lyft_rnd_res,
                    lin_lyft_bayes_res,
                    tree_lyft_cv,
                    tree_lyft_rnd_res,
                    tree_lyft_bayes_res,
                    elastic_net_lyft_cv,
                    elastic_net_lyft_rnd_res,
                    elastic_net_lyft_bayes_search,
                    xgb_lyft_grid,
                    xgb_lyft_Random,
                    xgb_lyft_bayes]
# Identify the best model
for df in dataframes_lyft:
    df.reset_index(drop=True, inplace=True)

subset_values_lyft = [df.loc[0, 'mean_test_score'] for df in dataframes_lyft]

max_value_lyft = max(subset_values_lyft)
dataframe_name = [name for name, obj in globals().items() if obj is_
↳dataframes_lyft[subset_values_lyft.index(max_value_lyft)]] [0]

print("Best Model for Lyft prediction is", dataframe_name, ": ", max_value_lyft)

```

Best Model for Uber prediction is xgb_uber_bayes : -1.886856033451432
Best Model for Lyft prediction is xgb_lyft_bayes : -2.9648903786512717

```

[ ]: # Uber actual vs. predicted values
y_pred_bayes_uber = xgb_uber_bayes_search.best_estimator_.predict(X_test_uber)

result_df_uber = pd.DataFrame({'Actual': y_test_uber, 'Predicted':_
↳y_pred_bayes_uber})
result_df_uber = result_df_uber.sort_values(by='Actual')


[ ]: # Plot a scatter plot to compare actual vs. predicted values
plt.figure(figsize=(8, 6))

```

```

sns.scatterplot(x='Actual', y='Predicted', data=result_df_uber,
               palette='viridis')
plt.title('Actual vs. Predicted Values for Uber')
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.show()

```



```

[ ]: import seaborn as sns
import matplotlib.pyplot as plt


result_df_uber['diff'] = abs(result_df_uber['Actual'] -
                             result_df_uber['Predicted']).round()

sorted_diff_uber = result_df_uber['diff'].sort_values(ascending=True)

[ ]: # Create a count plot using seaborn
plt.figure(figsize=(16, 4))
sns.countplot(x = sorted_diff_uber)
plt.title('Absolute Differences between Actual and Predicted Values')

```

```
plt.xlabel('Absolute Difference')
plt.ylabel('Count')
plt.show()
```



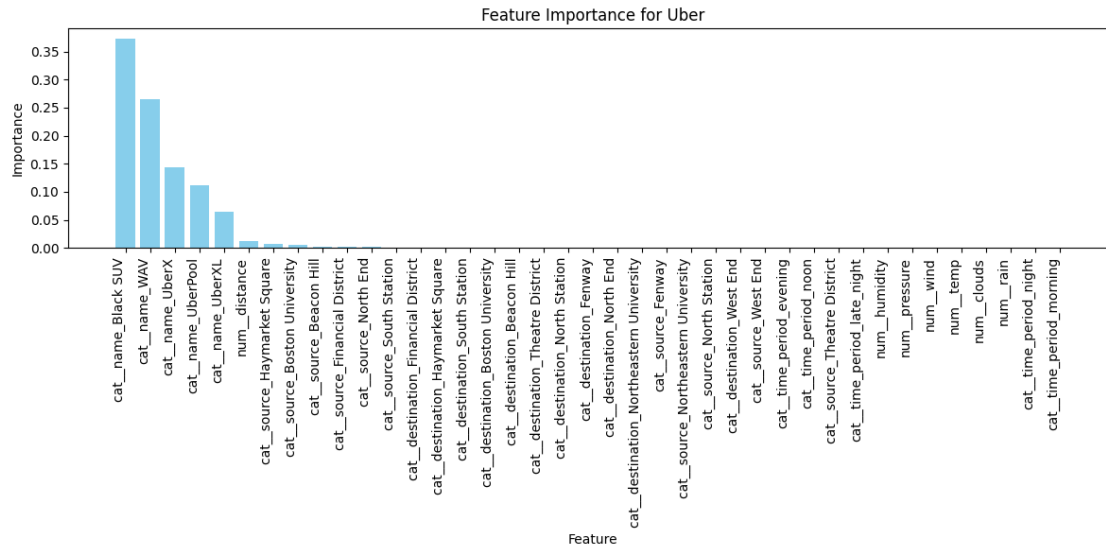
```
[ ]: # plotting importance feature for Uber
feature_importances_uber = xgb_uber_bayes_search.best_estimator_.
    ↪named_steps['xgbregressor'].feature_importances_

feature_names_uber = xgb_uber_bayes_search.best_estimator_[:-1].
    ↪get_feature_names_out()

importance_df_uber = pd.DataFrame({'Feature': feature_names_uber, 'Importance':
    ↪feature_importances_uber})

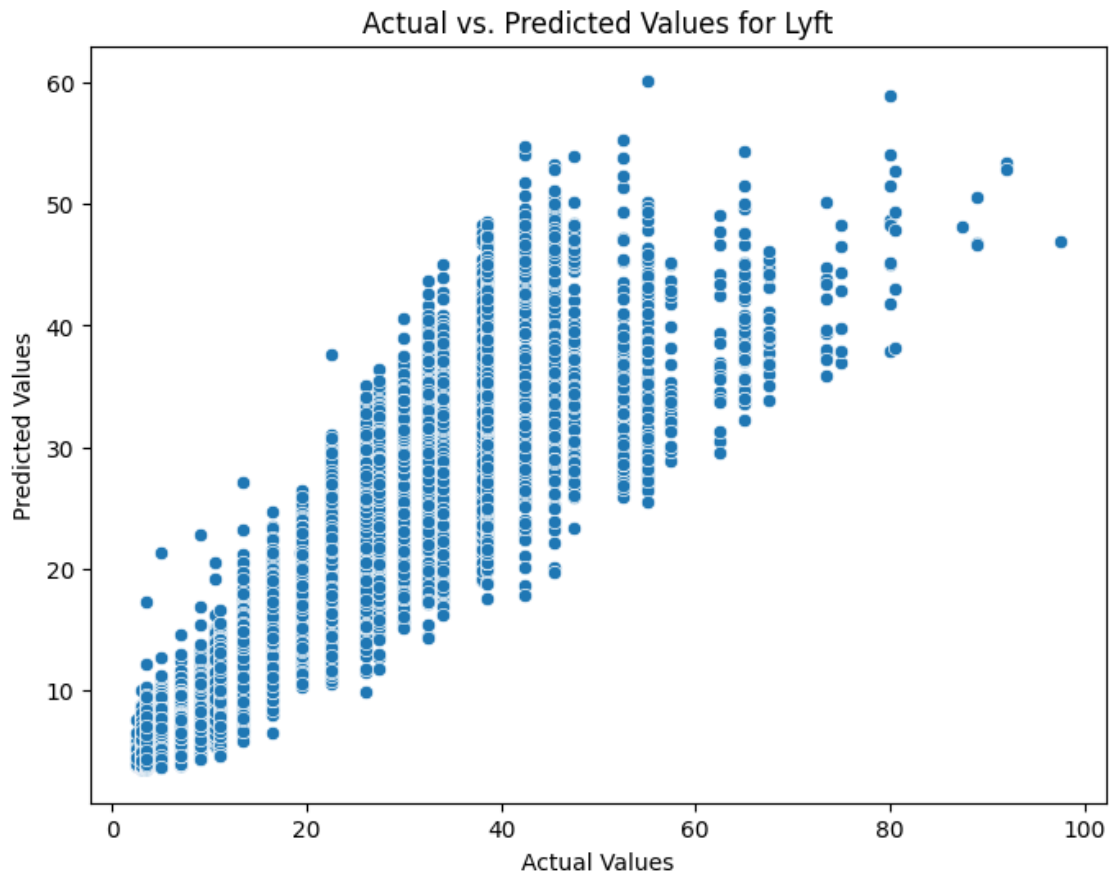
importance_df_uber = importance_df_uber.sort_values(by='Importance',
    ↪ascending=False)

plt.figure(figsize=(12, 6))
plt.bar(importance_df_uber['Feature'], importance_df_uber['Importance'],
    ↪color='skyblue')
plt.xlabel('Feature')
plt.ylabel('Importance')
plt.title('Feature Importance for Uber')
plt.xticks(rotation=90, ha='right')
plt.tight_layout()
plt.show()
```



```
[ ]: # Lyft actual vs. predicted values
y_pred_tree_lyft = xgb_lyft_bayes_search.predict(X_test_lyft)
result_df_lyft = pd.DataFrame({'Actual': y_test_lyft, 'Predicted':
    ↳ y_pred_tree_lyft})
result_df_lyft = result_df_lyft.sort_values(by='Actual')
```


```
[ ]: # plotting compare actual against predicted values
plt.figure(figsize=(8, 6))
sns.scatterplot(x='Actual', y='Predicted', data=result_df_lyft,
    ↳ palette='viridis')
plt.title('Actual vs. Predicted Values for Lyft')
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.show()
```



```
[ ]: result_df_lyft['diff'] = abs(result_df_lyft['Actual'] -
    ↪ result_df_lyft['Predicted']).round()

sorted_diff_lyft = result_df_lyft['diff'].sort_values(ascending=True)
```

```
[ ]: # Create a count plot using seaborn
plt.figure(figsize=(16, 4))
sns.countplot(x=sorted_diff_lyft)
plt.title('Absolute Differences between Actual and Predicted Values')
plt.xlabel('Absolute Difference')
plt.ylabel('Count')
plt.show()
```



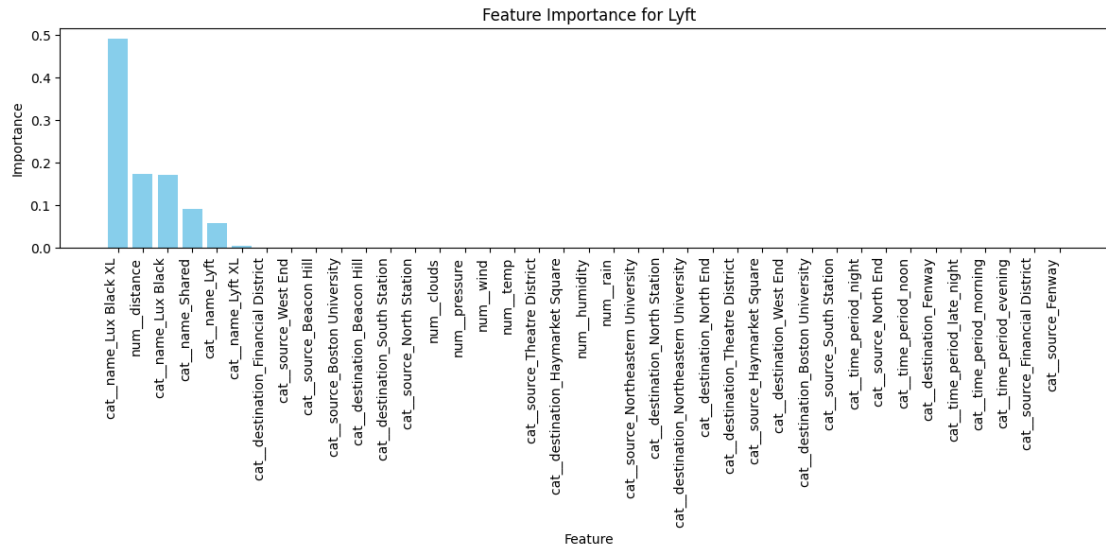
```
[ ]: # plotting importance feature for Lyft
feature_importances_lyft = tree_lyft_bayesian_search.best_estimator_.
    ↪named_steps['decisiontreeregressor'].feature_importances_

feature_names_lyft = tree_lyft_bayesian_search.best_estimator_[:-1].
    ↪get_feature_names_out()

importance_df_lyft = pd.DataFrame({'Feature': feature_names_lyft, 'Importance':
    ↪feature_importances_lyft})

importance_df_lyft = importance_df_lyft.sort_values(by='Importance',
    ↪ascending=False)

plt.figure(figsize=(12, 6))
plt.bar(importance_df_lyft['Feature'], importance_df_lyft['Importance'],
    ↪color='skyblue')
plt.xlabel('Feature')
plt.ylabel('Importance')
plt.title('Feature Importance for Lyft')
plt.xticks(rotation=90, ha='right')
plt.tight_layout()
plt.show()
```

6 Price prediction

```
[ ]: # Price prediction for Uber
test_data_uber = pd.DataFrame({
    'distance': [4.05],
    'destination': 'Financial District',
    'source': 'Boston University',
    'cab_type' : ['Uber'],
    'name': ['UberXL'],
    'temp': [39],
    'clouds': [0.7],
    'pressure': [1008],
    'rain': [0.05],
    'humidity': [0.76],
    'wind': [6.8],
    'time_period': ['afternoon']
})
test_data_uber
```

```
[ ]:   distance      destination      source cab_type   name  temp \
0      4.05  Financial District  Boston University    Uber  UberXL   39

   clouds  pressure  rain  humidity  wind  time_period
0      0.7     1008  0.05      0.76   6.8    afternoon
```

```
[ ]: # Price prediction for Lyft
test_data_lyft = pd.DataFrame({
```

```

    'distance': [4.05],
    'destination': 'Financial District',
    'source': 'Boston University',
    'cab_type' : ['Lyft'],
    'name': ['Lyft XL'],
    'temp': [39],
    'clouds': [0.7],
    'pressure': [1008],
    'rain': [0.05],
    'humidity': [0.76],
    'wind': [6.8],
    'time_period': ['afternoon']
})
test_data_lyft

```

```

[ ]: distance destination source cab_type name temp \
0      4.05 Financial District Boston University Lyft Lyft XL 39

      clouds pressure rain humidity wind time_period
0      0.7      1008 0.05      0.76 6.8 afternoon

```

```

[ ]: data_predictions_uber = xgb_uber_bayes_search.best_estimator_.
      ↪predict(test_data_uber)
print("The prediction price of Uber is", data_predictions_uber)
data_predictions_lyft = xgb_lyft_bayes_search.best_estimator_.
      ↪predict(test_data_lyft)
print("The prediction price of Lyft is", data_predictions_lyft)

```

The prediction price of Uber is [23.334993]
The prediction price of Lyft is [22.248138]

```

[ ]: # Change the destination
test_data_uber_1 = test_data_uber.copy()
test_data_lyft_1 = test_data_lyft.copy()
test_data_uber_1['destination'] = 'Theatre District'
test_data_lyft_1['destination'] = 'Theatre District'
test_data_uber_1['distance'] = 2.9
test_data_lyft_1['distance'] = 2.9
data_predictions_uber = xgb_uber_bayes_search.best_estimator_.
      ↪predict(test_data_uber_1)
print("The prediction price of Uber is", data_predictions_uber)
data_predictions_lyft = xgb_lyft_bayes_search.best_estimator_.
      ↪predict(test_data_lyft_1)
print("The prediction price of Lyft is", data_predictions_lyft)

```

The prediction price of Uber is [18.100248]
The prediction price of Lyft is [20.249187]

```
[ ]: # Change the time_period
test_data_uber_2 = test_data_uber.copy()
test_data_lyft_2 = test_data_lyft.copy()
test_data_uber_2['time_period'] = 'evening'
test_data_lyft_2['time_period'] = 'evening'
data_predictions_uber = xgb_uber_bayes_search.best_estimator_.
    ↪predict(test_data_uber_2)
print("The prediction price of Uber is", data_predictions_uber)
data_predictions_lyft = xgb_lyft_bayes_search.best_estimator_.
    ↪predict(test_data_lyft_2)
print("The prediction price of Lyft is", data_predictions_lyft)
```

The prediction price of Uber is [23.310923]
The prediction price of Lyft is [22.345005]

```
[ ]: # Change the Vehicle type
test_data_uber_2 = test_data_uber.copy()
test_data_lyft_2 = test_data_lyft.copy()
test_data_uber_2['name'] = 'Black SUV'
test_data_lyft_2['name'] = 'Lux Black XL'
data_predictions_uber = xgb_uber_bayes_search.best_estimator_.
    ↪predict(test_data_uber_2)
print("The prediction price of Uber is", data_predictions_uber)
data_predictions_lyft = xgb_lyft_bayes_search.best_estimator_.
    ↪predict(test_data_lyft_2)
print("The prediction price of Lyft is", data_predictions_lyft)

test_data_uber_1
```

The prediction price of Uber is [40.76876]
The prediction price of Lyft is [43.09772]

```
[ ]: distance      destination      source cab_type  name  temp  \
0         2.9  Theatre District  Boston University    Uber  UberXL    39

      clouds  pressure  rain  humidity  wind  time_period
0         0.7    1008  0.05     0.76   6.8    afternoon
```

7 Predictions on test data

```
[ ]: # Uber predictions with best estimator
data_predictions_uber = xgb_uber_bayes_search.best_estimator_.
    ↪predict(X_test_uber)
data_predictions_uber = pd.DataFrame(data_predictions_uber)
X_test_uber = X_test_uber.reset_index()
X_test_uber['predicted_price'] = data_predictions_uber
```

```
X_test_uber.head(10)
```

```
[ ]:      index  distance  cab_type      destination \
0  460051      3.01    Uber      Boston University
1  235852      0.56    Uber      Haymarket Square
2  521026      2.72    Uber      West End
3  369283      1.35    Uber      North End
4  406967      1.08    Uber  Northeastern University
5  657376      2.65    Uber      Theatre District
6  481690      1.35    Uber      Haymarket Square
7  399507      2.45    Uber      North Station
8  528035      2.34    Uber      Back Bay
9   51934      1.92    Uber  Northeastern University

      source      name      temp      clouds      pressure \
0      North Station      WAV  32.280000  0.690000  1033.670000
1      North Station  UberX  50.740000  1.000000  1003.250000
2      Fenway      UberX  40.430000  0.890000  1014.006000
3      Beacon Hill  UberXL  41.310000  0.893333  991.560000
4      Back Bay      Black  41.153333  0.633333  991.586667
5      Fenway      WAV  48.600000  0.560000  1021.350000
6      Beacon Hill  Black SUV  23.980000  0.500000  1008.790000
7  Northeastern University      WAV  19.940000  0.450000  1031.710000
8      North End  UberPool  42.960000  0.965000  988.545000
9      Beacon Hill  Black SUV  51.840000  0.710000  1021.600000

      rain  humidity      wind  time_period  predicted_price
0  0.000  0.560000  2.360000      evening      11.093656
1  0.000  0.930000  3.930000      evening      7.184306
2  0.000  0.932000  2.634000  late_night      10.275254
3  0.002  0.633333  9.490000      afternoon      13.635898
4  0.000  0.623333  10.113333      afternoon      15.626392
5  0.000  0.720000  5.370000      evening      10.958372
6  0.000  0.510000  14.930000      noon      26.742880
7  0.000  0.630000  2.740000      noon      11.089293
8  0.000  0.895000  9.720000      afternoon      10.013367
9  0.000  0.750000  6.740000      afternoon      26.821289
```

```
[ ]: # Lyft predictions with best estimator
data_predictions_lyft = xgb_lyft_bayes_search.best_estimator_.
    predict(X_test_lyft)
data_predictions_lyft = pd.DataFrame(data_predictions_lyft)
X_test_lyft = X_test_lyft.reset_index()
X_test_lyft['predicted_price'] = data_predictions_lyft
X_test_lyft.head(10)
```

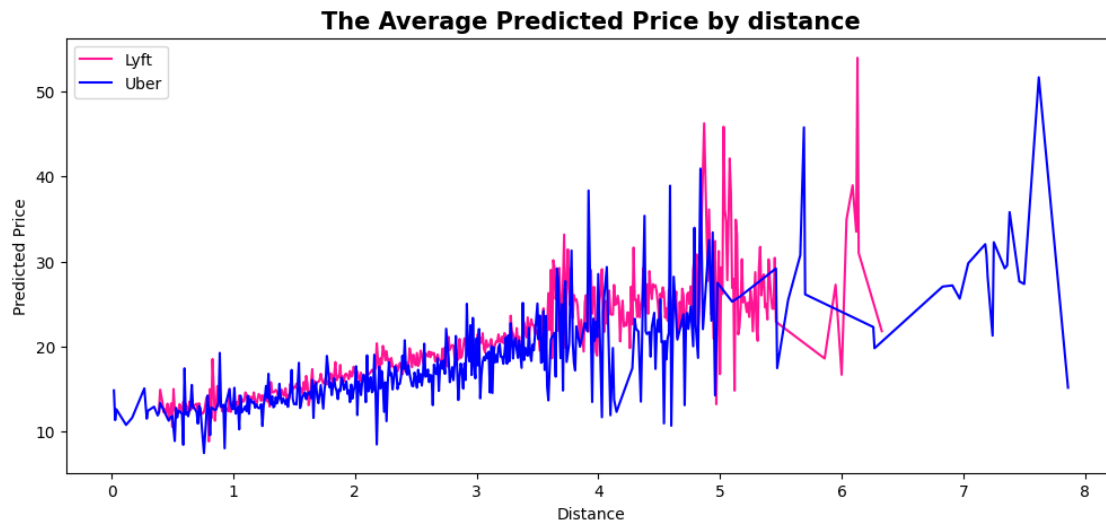
```
[ ]:      index  distance cab_type      destination      source \
0  635446      3.19      Lyft  Boston University  North Station
1   13360      3.16      Lyft                Fenway  North Station
2  368367      2.97      Lyft                Fenway    West End
3  324716      0.74      Lyft  Haymarket Square    West End
4  574640      3.15      Lyft  Boston University  North Station
5  496732      2.86      Lyft  Boston University    West End
6  553575      0.53      Lyft    North Station  Haymarket Square
7  215361      2.39      Lyft    North Station    South Station
8  558301      3.17      Lyft  Boston University  Theatre District
9  122649      3.10      Lyft                Fenway  Theatre District

      name      temp      clouds  pressure      rain  humidity      wind \
0    Shared  46.920  0.910000  1022.170  0.0000  0.880000  3.990000
1      Lyft  45.060  0.980000  1012.140  0.0000  0.880000  5.160000
2  Lux Black XL  39.210  0.440000  1022.810  0.0000  0.740000  7.190000
3    Shared  37.940  1.000000  1020.070  0.0183  0.860000  1.920000
4    Shared  29.880  0.690000  1034.970  0.0000  0.680000  1.950000
5  Lux Black XL  42.860  0.675000   991.295  0.0000  0.715000  8.565000
6  Lux Black XL  42.595  0.790000   990.810  0.0000  0.725000  9.325000
7  Lux Black XL  33.575  0.138333   991.210  0.0000  0.821667  5.773333
8    Lyft XL  29.820  0.710000  1035.070  0.0000  0.720000  1.610000
9    Lyft    34.100  1.000000  1003.700  0.0100  0.940000  9.320000

time_period  predicted_price
0  afternoon           6.929302
1   evening          10.729082
2   morning          33.897461
3   morning           4.331971
4    night           6.764403
5   evening          34.063202
6   evening          27.014526
7     noon          32.993515
8  late_night          18.120478
9     noon          12.647260
```

```
[ ]: # plotting distance against predicted price
fig , ax = plt.subplots(figsize = (12,5))
ax.plot(X_test_lyft.groupby('distance').predicted_price.mean().index,
        ↪X_test_lyft.groupby('distance').predicted_price.mean(), label = 'Lyft',
        ↪color='deeppink')
ax.plot(X_test_uber.groupby('distance').predicted_price.mean().index,
        ↪X_test_uber.groupby('distance').predicted_price.mean(), label = 'Uber',
        ↪color='blue')
ax.set_title('The Average Predicted Price by distance', fontsize= 15,
        ↪fontweight='bold')
ax.set(xlabel = 'Distance', ylabel = 'Predicted Price' )
```

```
ax.legend()
plt.show()
```



Upon comparing the actual price with the average predicted price, we observed that the latter is consistently higher. Furthermore, beyond a distance of five miles, the average predicted price shows greater fluctuation, potentially attributed to a reduced number of data points.

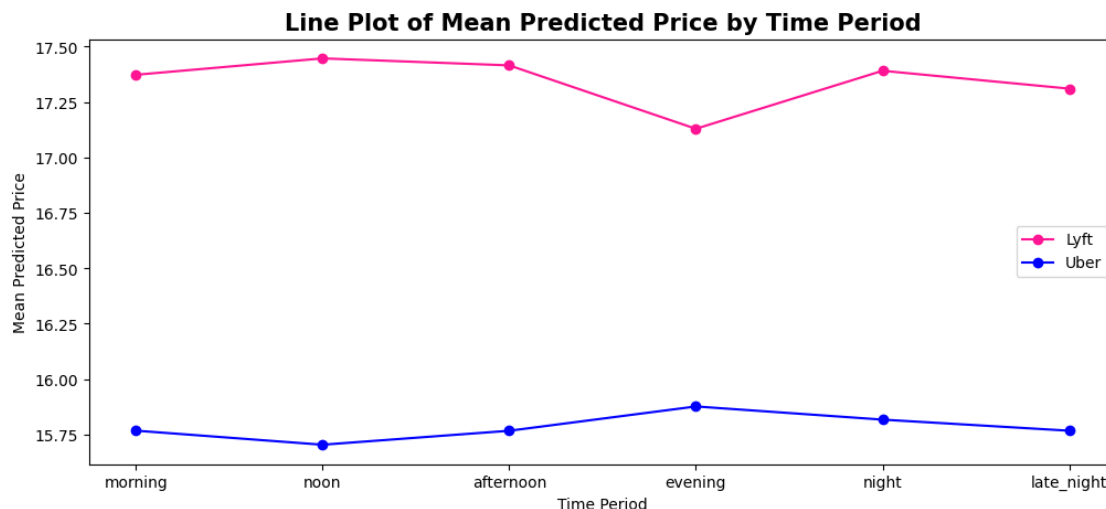
```
[ ]: # plotting time period against predicted price
time_period_order = ['morning', 'noon', 'afternoon', 'evening', 'night',
                    ↪ 'late_night']
fig, ax = plt.subplots(figsize=(12, 5))

# reorder the dataframe based on time_period_order
lyft_df_ordered = X_test_lyft.groupby('time_period').predicted_price.mean().
                    ↪ loc[time_period_order]
uber_df_ordered = X_test_uber.groupby('time_period').predicted_price.mean().
                    ↪ loc[time_period_order]

ax.plot(lyft_df_ordered.index, lyft_df_ordered, label='Lyft', color='deeppink',
        ↪ linestyle='-', marker='o')
ax.plot(uber_df_ordered.index, uber_df_ordered, label='Uber', color='blue',
        ↪ linestyle='-', marker='o')

ax.set_title('Line Plot of Mean Predicted Price by Time Period', fontsize=15,
            ↪ fontweight='bold')
ax.set(xlabel='Time Period', ylabel='Mean Predicted Price')
ax.legend()

plt.show()
```



Overall, the average predicted price for Lyft is constantly higher than Uber.

Unexpectedly, we observed a price decrease during the evening period for Lyft, in contrast to a price increase for Uber. We speculate that Lyft reduces prices during this time to enhance competitiveness, particularly during peak demand hours.

8 Challenges

- When joining the trips and weather tables, we found out that the data in the weather table were queried by a 5 minute frequency. The weather table could not directly be joined on the trips table timestamp, since the date-time might not be exactly the same. To overcome this challenge, we added a new column to take the average weather data by hour, and matched it to the trips table on the hour extracted from the timestamp.
- We had to pick two best models for the Uber and Lyft dataset separately, and since each model had 4 models with 3 search methods, it was challenging to efficiently extract the models with the best performing hyperparameters. To solve this problem, we used for loops to iterate through the different models and the mean test scores. We then used list comprehensions to extract the model with the maximum mean test score and printed out the name of the best-performing models.
- With a large dataset, we had to limit the number of hyperparameters that we tune in order to run the models efficiently.

9 Conclusion

In conclusion, the XGBoost regression model demonstrated the best performance for predicting both Uber and Lyft prices, with the lowest root mean squared error scores compared to other models tested.

By conducting supervised machine learning to predict Uber and Lyft surge prices, we learned the

importance of finding the most accurate and cost-effective machine learning model to solve business problems with limited resources. With time and compute power limitations, we needed to use a select number of hyperparameters to tune our models and find the best prediction model for Uber and Lyft separately.

Feature importance analysis found that the distance traveled was the most influential factor affecting price predictions for both Uber and Lyft. This aligns with intuition as trip distance directly impacts the fare amount. Other top predictors were location attributes like pick-up location, destination, and time period.

Comparisons of actual versus predicted price values showed the XGBoost models were relatively accurate, with most predictions closely aligned or slightly under-estimating the actual fares. The predicted prices fell within a \$5 difference for a majority of the test cases.

Overall, our analysis and modeling provides an effective framework to predict Uber and Lyft pricing based on a variety of features like trip details, geographic locations, vehicle types and external factors like weather. The XGBoost approach outperformed other regression models, demonstrating feasibility to deliver accurate price estimates to riders using historical Uber and Lyft data trends. With some refinement, the predictive models show promise to be integrated into passenger travel planning and budgeting use cases going forward.